

Memory Virtualization

Questions answered in this lecture:

What is in the address space of a process (review)?

What are the different ways that that OS can virtualize memory?


Time sharing, static relocation, dynamic relocation

(base, base + bounds, segmentation)

What hardware support is needed for dynamic relocation?

Motivation for Virtualization

- **Uniprogramming: One process runs at a time**



- **Disadvantages:**
 - Only one process runs at a time
 - Process can destroy OS

Multiprogramming Goals

■ Transparency

- Processes are **not aware** that memory is shared
- Works regardless of number and/or location of processes

■ Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

■ Efficiency


- Do not waste memory resources (minimize fragmentation)

■ Sharing

- Cooperating processes can share portions of address space

Abstraction: Address Space

- **Address space**: Each process has set of addresses that map to bytes
- **Review**: What is **in** an address space?
- **Address space has static and dynamic components**
 - **Static**: Code and some global variables
 - **Dynamic**: Stack and Heap



Motivation for Dynamic Memory

■ Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
 - Allocate enough for worst possible case; Storage is used inefficiently

■ Recursive procedures

- Do not know how many times procedure will be nested

■ Complex data structures: lists and trees

- ```
struct my_t *p = (struct my_t *)
 malloc(sizeof(struct my_t));
```

## ■ Two types of dynamic allocation

- Stack
- Heap

# Stack Organization

- **Definition: Memory is freed in opposite order from allocation**

```
alloc(A);
alloc(B);
alloc(C);
free(C);
alloc(D);
free(D);
free(B);
free(A);
```

- **Simple and efficient implementation:**
  - Pointer separates allocated and freed space**
- Allocate: Increment pointer
- Free: Decrement pointer
- **No fragmentation**

# Where Are Stacks Used?

- OS uses stack for procedure call **frames** (local variables and parameters)

```
main () {
 int A = 0;
 foo (A);
 printf("A: %d\n", A);
}
```

```
void foo (int Z) {
 int A = 2;
 Z = 5;
 printf("A: %d Z: %d\n", A, Z);
}
```

# Heap Organization

- **Definition: Allocate from any random location: malloc(), new()**
  - Heap memory consists of allocated areas and free areas (holes)
  - Order of allocation and free is unpredictable

- **Advantage**

- Works for all data structures

- **Disadvantages**

- Allocation can be slow
  - End up with small chunks of free space - fragmentation
  - Where to allocate 12 bytes? 16 bytes? 24 bytes??

- **What is OS's role in managing heap?**

- OS gives big chunk of free memory to process; library manages individual allocations





# x86-64 Linux Memory Layout

*not drawn to scale*

## ■ Stack

- Runtime stack (8MB limit)
- E. g., local variables

## ■ Heap

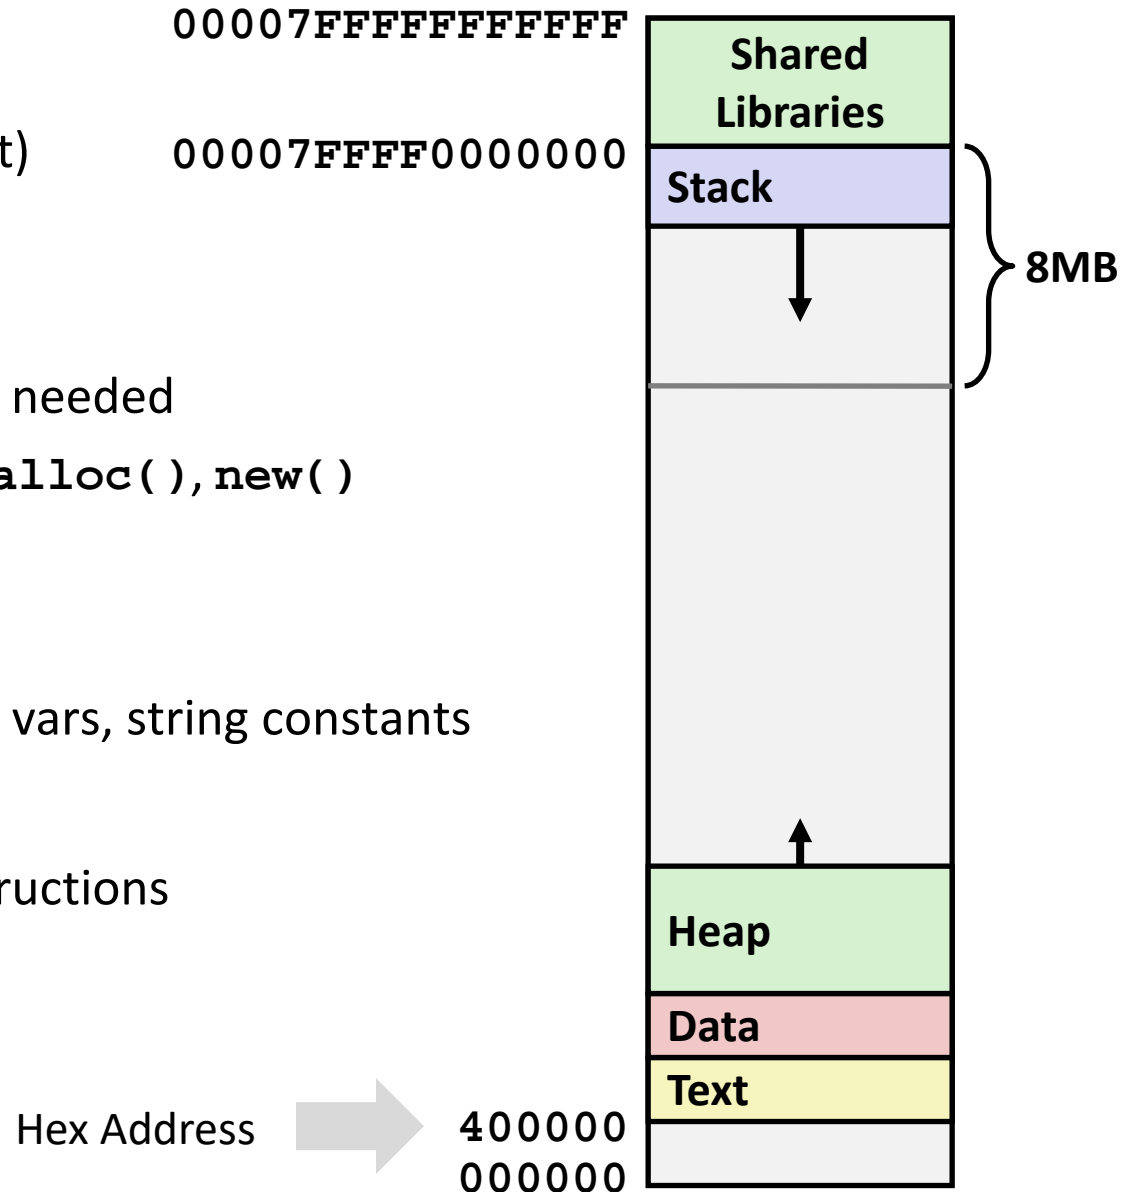
- **Dynamically** allocated as needed
- When call `malloc()`, `calloc()`, `new()`

## ■ Data

- **Statically** allocated data
- E.g., global vars, `static` vars, string constants

## ■ Text / Shared Libraries

- Executable machine instructions
- Read-only



# Memory Access

```
int function(int len, char* src) {
 char buffer[100];
 process(buffer, src, len);
 return 0;
}
```

```
void main() {
 function(50, "Hello\n");
}
```

Calling Convention:  
function(%rdi, %rsi, %rdx)


```
sub $0x64, %rsp
movq %rdi, %rdx
movq %rsp, %rdi
callq < process >
xor %rax,%rax
add $0x64, %rsp
retq
```


# How to Virtualize Memory?

- Problem: How to run multiple processes simultaneously?
- Addresses are “hardcoded” into process binaries
- How to avoid collisions?
- Possible Solutions for Mechanisms (covered today):
  1. Time Sharing
  2. Static Relocation
  3. Base
  4. Base+Bounds
  5. Segmentation

# 1) Time Sharing of Memory

- Try similar approach to how OS virtualizes CPU
- **Observation:**  
OS gives illusion of many virtual CPUs by saving CPU registers to memory when a process isn't running
- Could give illusion of many virtual memories by saving memory to disk when process isn't running





# Problems with Time Sharing Memory

- **Problem**

- Ridiculously poor performance

- **Better Alternative: space sharing**

- At same time, space of memory is divided across processes

- **Remainder of solutions all use space sharing**

## 2) Static Relocation

- Idea: OS **rewrites** each program before **loading** it as a process in memory
- Each rewrite for different process uses **different addresses and pointers**
- Change jumps, loads of static data

- 0x10: movl 0x8(%rbp), %edi
- 0x13: addl \$0x3, %edi
- 0x19: movl %edi, 0x8(%rbp)

rewrite

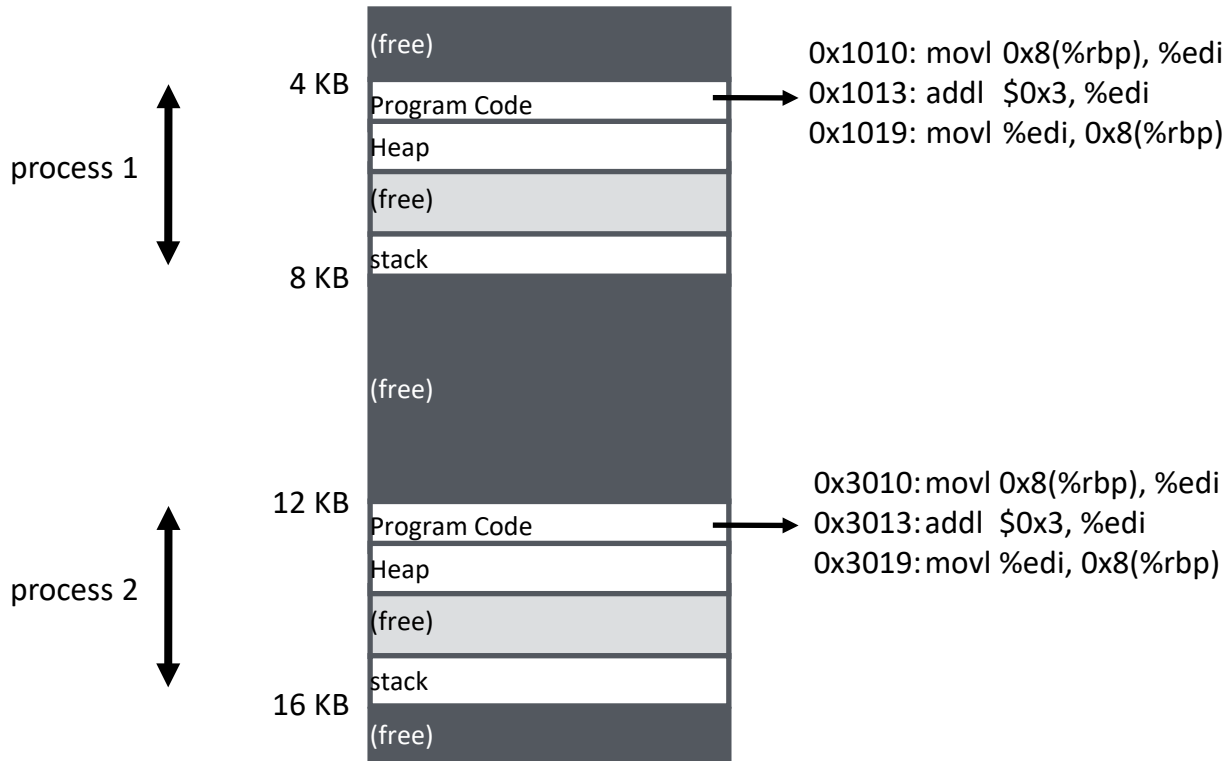
0x1010: movl 0x8(%rbp), %edi  
0x1013: addl \$0x3, %edi  
0x1019: movl %edi, 0x8(%rbp)

rewrite

0x3010: movl 0x8(%rbp), %edi  
0x3013: addl \$0x3, %edi  
0x3019: movl %edi, 0x8(%rbp)



# Static: Layout in Memory



# Static Relocation: Disadvantages

## ■ No protection

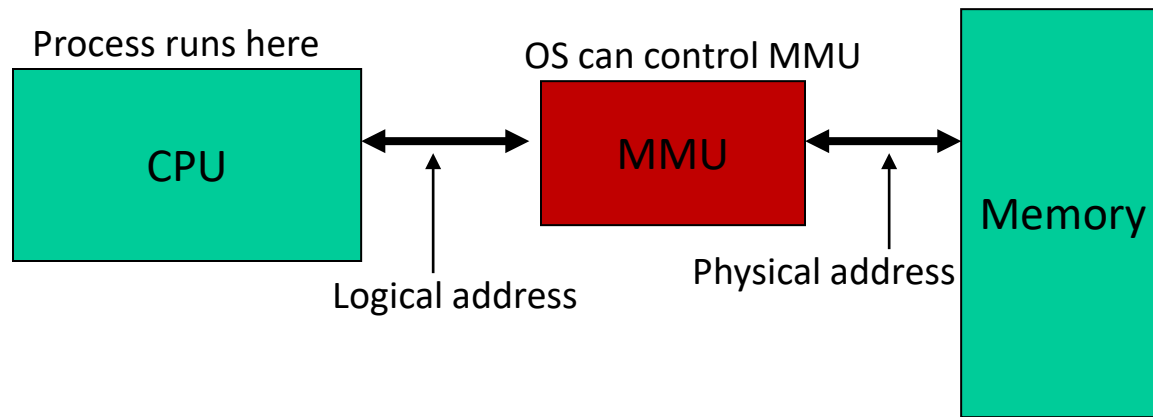
- Process can destroy OS or other processes
- No privacy

## ■ Cannot move address space after it has been placed

- May not be able to allocate new process
- Need to be the same place for being switched back

# 3) Dynamic Relocation

- Goal: Protect processes from one another
- Requires hardware support
  - Memory Management Unit (MMU)
- MMU dynamically changes process address at every memory reference
  - Process generates **logical** or **virtual** addresses (in their address space)
  - Memory hardware uses **physical** or **real** addresses



# Hardware Support for Dynamic Relocation

## ■ Two operating modes

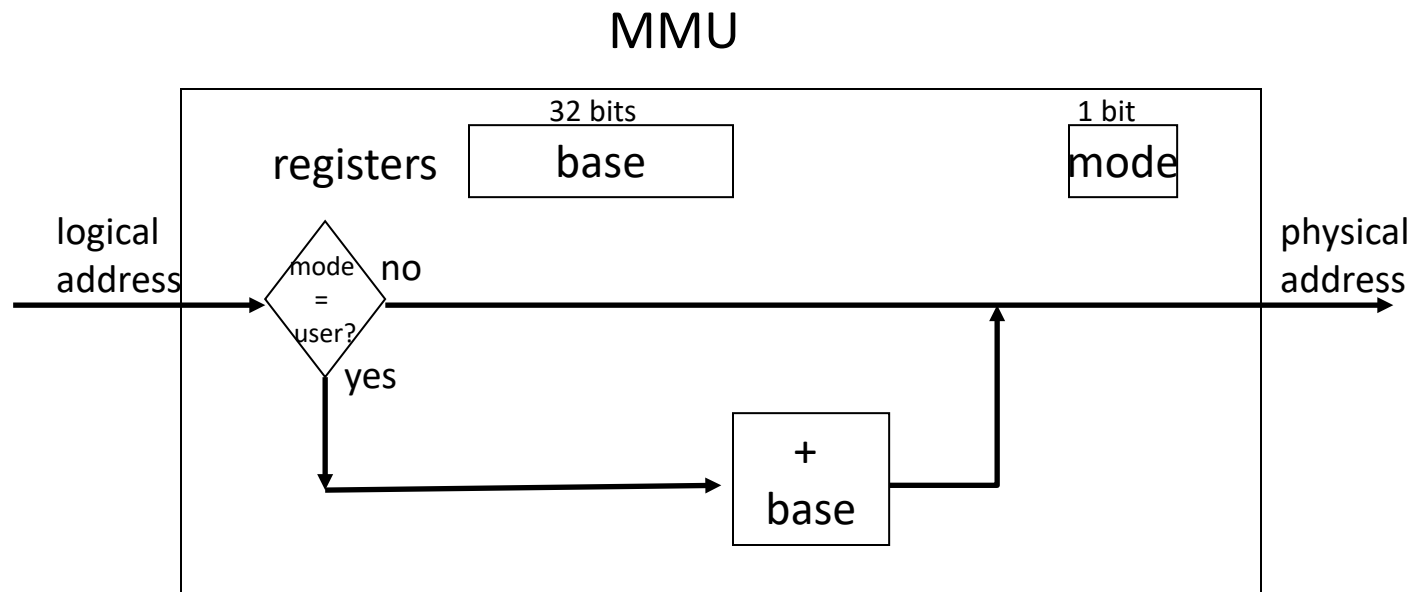
- **Privileged (protected, kernel) mode: OS runs**
  - When enter **OS** (trap, system calls, interrupts, exceptions)
  - Allows **certain instructions** to be executed
    - Can manipulate contents of MMU
  - Allows OS to access **all of physical memory**
- **User mode: User processes run**
  - Perform translation of logical address to physical address

## ■ Minimal MMU contains base register for translation

- base: start location for address space


# Implementation of Dynamic Relocation: BASE REG

- Translation on **every memory access** of user process
  - MMU adds base register to logical address to form physical address




# Dynamic Relocation with Base Register

- Idea: translate virtual addresses to physical by **adding a fixed offset** each time.
- Store **offset** in **base register**
- **Each process** has **different** value in base register




P1 is running




P2 is running






| Virtual            | Physical       |              |
|--------------------|----------------|--------------|
| P1: mov [100], R10 | load 1124, R10 | (1024 + 100) |



| Virtual            | Physical       |
|--------------------|----------------|
| P1: mov [100], R10 | load 1124, R10 |
| P2: mov [100], R10 | load 4196, R10 |



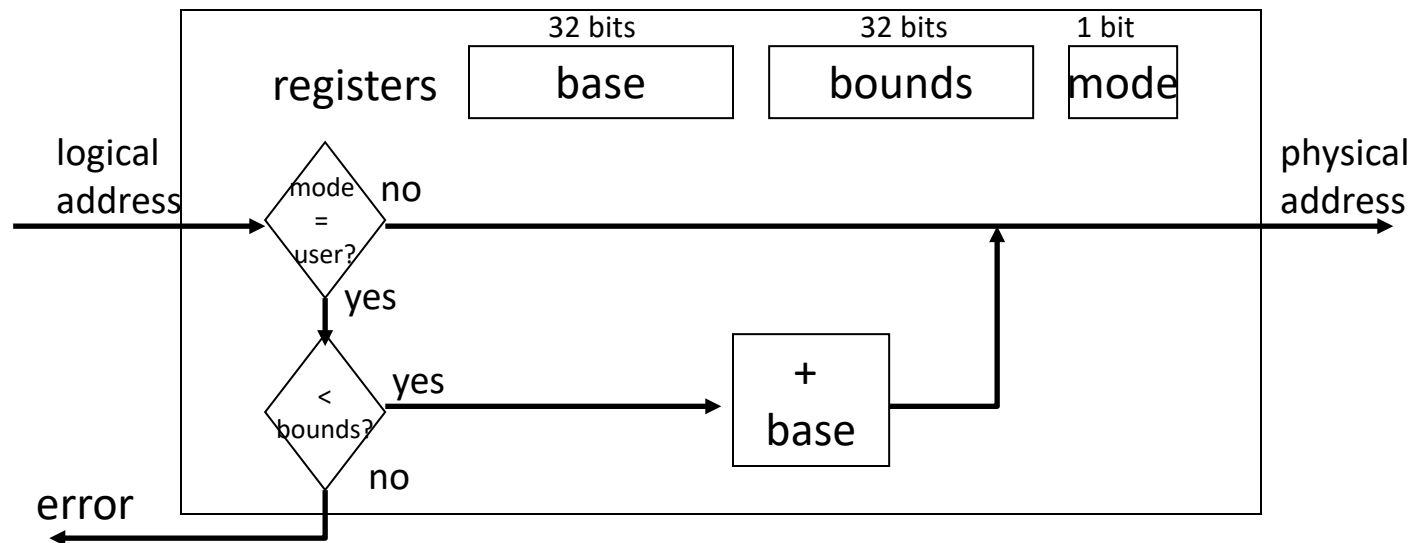
How well does dynamic relocation do with base register for protection?

## 4) Dynamic with Base+Bounds


- Idea: **limit** the address space with a bounds register
- **Base register**: smallest physical addr (or starting location)
- **Bounds register**: size of this process's virtual address space
  - Sometimes defined as largest physical address (base + size)
- OS kills process if process loads/stores beyond bounds

# Implementation of BASE+BOUNDS

- Translation on **every memory access** of user process
  - MMU compares logical address to **bounds register**
    - if logical address is greater, then generate error
  - MMU adds base register to logical address to form physical address







P2 is running

# Managing Processes with Base and Bounds

## ■ Context-switch

- Add base and bounds registers to **PCB**
- Steps
  - Change to **privileged mode**
  - **Save** base and bounds registers of old process
  - **Load** base and bounds registers of new process
  - Change to **user mode** and jump to new process

## ■ **Protection** requirement

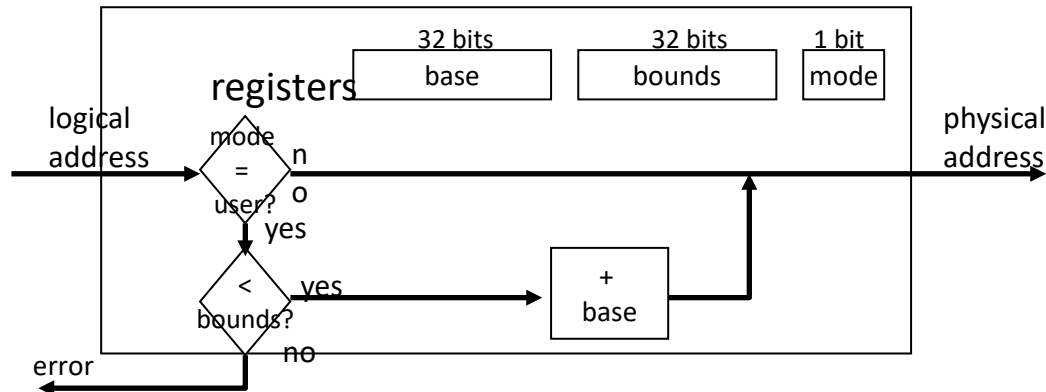
- User process **cannot** change base and bounds registers
- User process **cannot** change to privileged mode



# Base and Bounds Advantages

## ■ Advantages


- Provides **protection** (both read and write) across address spaces
- Supports **dynamic relocation**
  - Can place process at different locations initially and also move address spaces
- **Simple**, inexpensive implementation
  - Few registers, little logic in MMU
- **Fast**
  - Add and compare in parallel



# Base and Bounds DISADVANTAGES


## ■ Disadvantages

- Each process must be **allocated contiguously** in physical memory
  - Must allocate memory that may not be used by process
- **No partial sharing**: Cannot share limited parts of address space



# 5) Segmentation

- **Divide** address space into **logical segments**
  - Each segment corresponds to logical entity in address space
    - code, stack, heap
- Each segment can **independently**:
  - be placed separately in physical memory
  - grow and shrink
  - be protected (separate read/write/execute protection bits)



# Segmented Addressing

- Process now specifies **base and offset** within segment
- **How does process designate a particular segment?**
  - Use part of logical address
    - Top bits of logical address select segment
    - Low bits of logical address select offset within segment
- **What if small address space, not enough bits?**
  - Implicitly by type of memory reference
  - Special registers


# Segmentation Implementation

- **MMU contains Segment Table (per process)**
  - Each segment has own base and bounds, protection bits
  - **Example:** 14 bit logical address, 4 segments; how many bits for segment?  
How many bits for offset?

| Segment | Base   | Bounds | R W |
|---------|--------|--------|-----|
| 0       | 0x2000 | 0x6fff | 1 0 |
| 1       | 0x0000 | 0x4fff | 1 1 |
| 2       | 0x3000 | 0xffff | 1 1 |
| 3       | 0x0000 | 0x0000 | 0 0 |

remember:  
1 hex digit->4 bits

# Segmentation Implementation



```
Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
```

```
Offset = VirtualAddress & OFFSET_MASK
```

```
if (Offset >= Bounds[Segment])
```

```
 RaiseException(PROTECTION_FAULT)
```


```
else
```

```
 PhysAddr = Base[Segment] + Offset
```

```
 Register = AccessMemory(PhysAddr)
```

# Advantages of Segmentation

- **Enables sparse allocation of address space**
  - Stack and heap can grow independently
  - Heap: If no data on free list, dynamic memory allocator requests more from OS (e.g., UNIX: malloc calls sbrk())
  - Stack: OS recognizes reference outside legal segment, extends stack implicitly
- **Different protection for different segments**
  - Read-only status for code
- **Enables sharing of selected segments**
- **Supports dynamic relocation of each segment**



# Disadvantages of Segmentation

- Each segment must be allocated **contiguously**
  - May not have sufficient physical memory for large segments
- Fix in next lecture with paging...