

Virtualizing Memory: Paging

Questions answered in this lecture:

Review segmentation and fragmentation

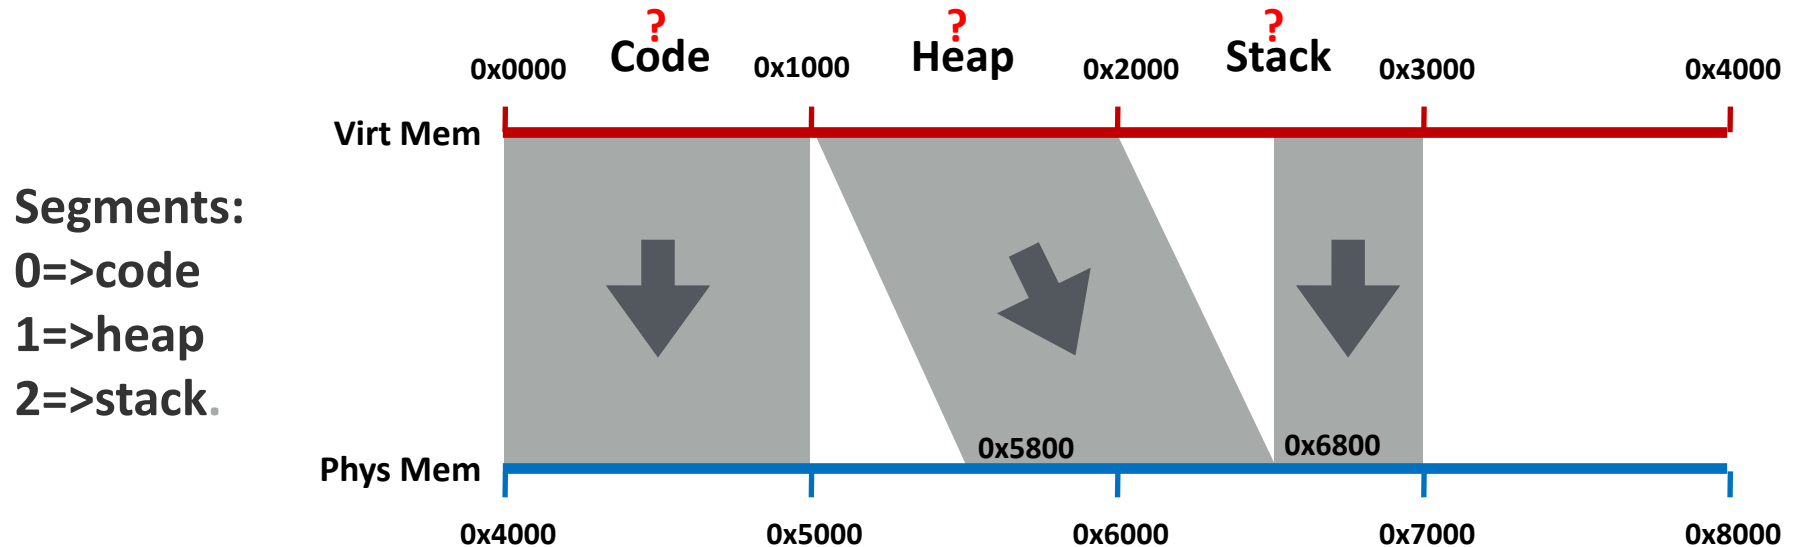
What is paging?

Where are page tables stored?

What are advantages and disadvantages of paging?

Review: Segmentation

Assume **14-bit** virtual addresses, high 2 bits indicate segment



Where doe segment table live?

All registers, MMU

Seg	Base	Bounds
0	0x4000	0xfff
1	0x5800	0xfff
2	0x6800	0x7ff

Review: Memory Accesses

```
0x0010: movl  (0x1100), %edi
0x0013: addl  $0x3, %edi
0x0019: movl  %edi, (0x1100)
```

%rip: 0x0010

Physical Memory Accesses?

1) Fetch instruction at logical addr 0x0010

- Physical addr: **0x4010**

Exec, load from logical addr 0x1100

- Physical addr: **0x5900**

2) Fetch instruction at logical addr 0x0013

- Physical addr: **0x4013**

Exec, no load

3) Fetch instruction at logical addr 0x0019

- Physical addr: **0x4019**

Exec, store to logical addr 0x1100

- Physical addr: **0x5900**

Total of 5 memory references (3 instruction fetches, 2 movl)

x86 Segmentation Implementation

■ Real Mode

segment selector
(16 bit register)



func:

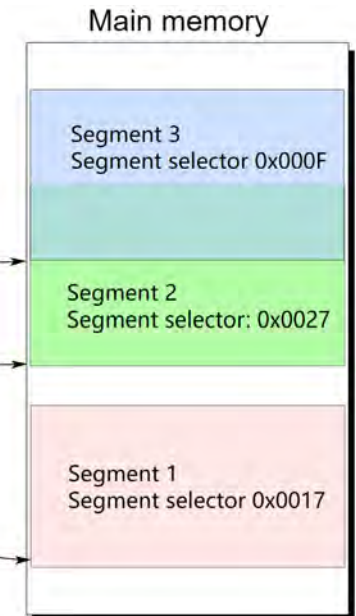
```
mov eax, [ss:esp+4]
mov ebx, [ds:eax]
add ebx, 8
mov [ds:eax], ebx
ret
```

0000 0110 1110 1111 0000	Segment	16 bits, shifted 4 bits left (or multiplied by 0x10)
+ 0001 0010 0011 0100	Offset	16 bits
<hr/>		
0000 1000 0001 0010 0100	Address	20 bits

■ Protected Mode

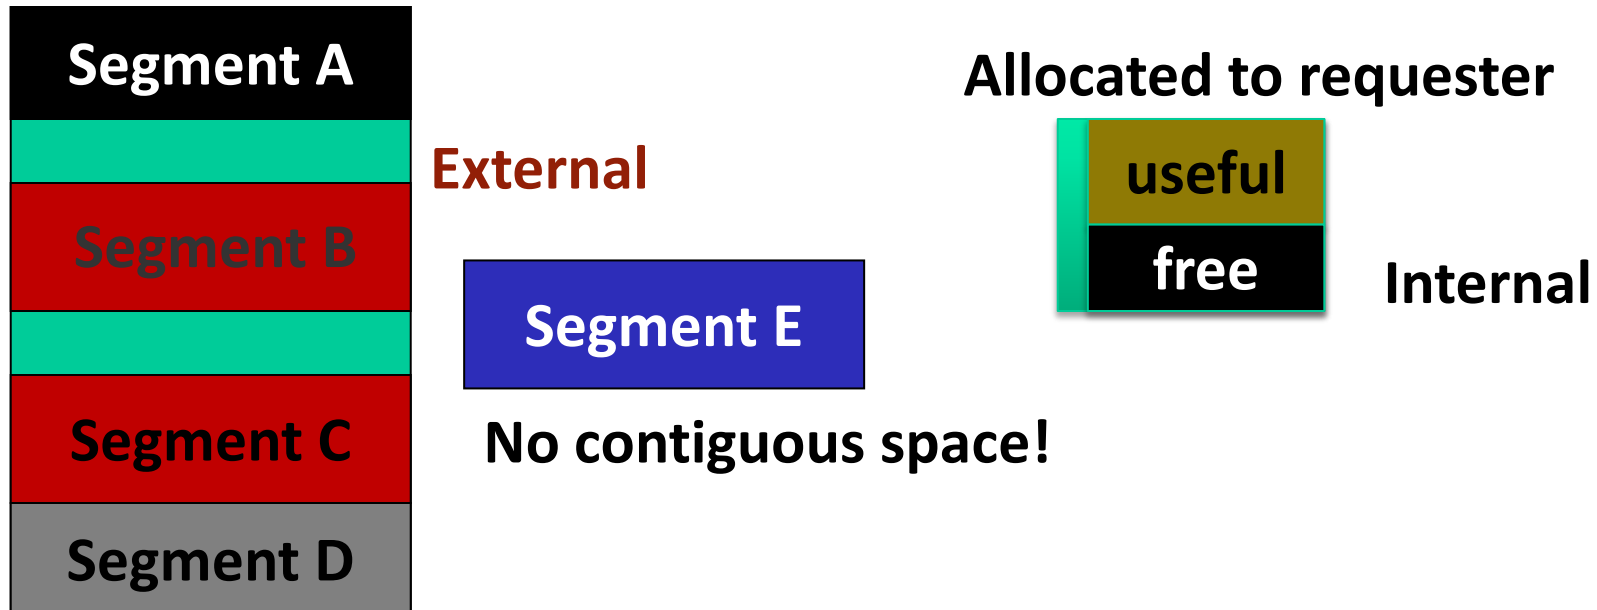
Local Descriptor Table (LDT)

5			
4	0x21430	0xC000	•
3			
2	0x0CEf0	0xA300	•
1	0x28C00	0xFC00	•
0			
Linear base address (BASE)		Segment size (LIMIT)	



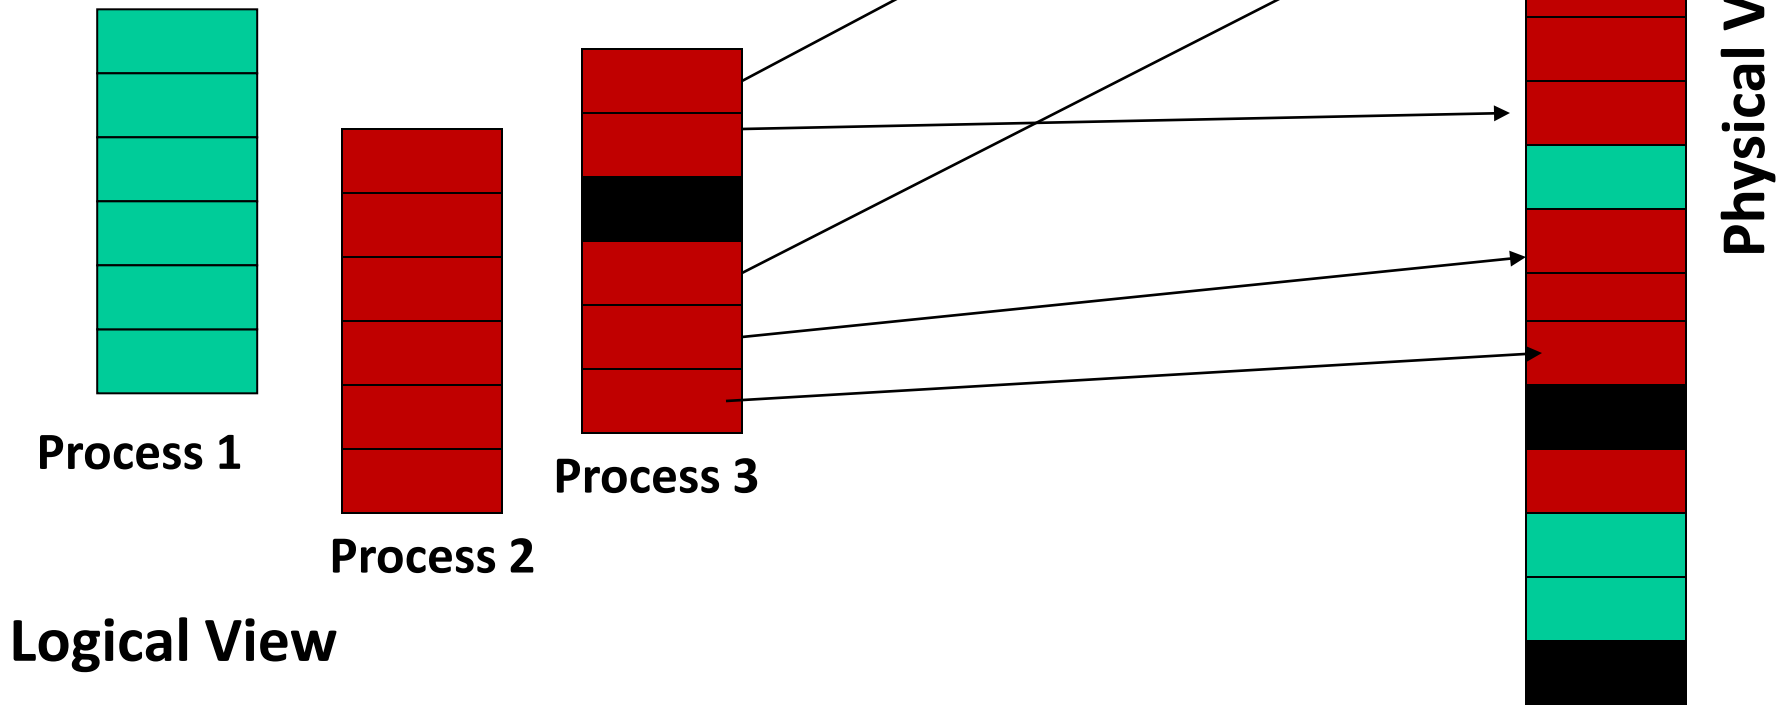
Problem: Fragmentation

- **Definition:** Free memory that can't be usefully allocated
- **Why?**
 - Free memory (hole) is too small and scattered
 - Rules for allocating memory prohibit using this free space
- **Types of fragmentation**
 - **External:** Visible to allocator (e.g., OS)
 - **Internal:** Visible to requester (e.g., if must allocate at some granularity)



Paging

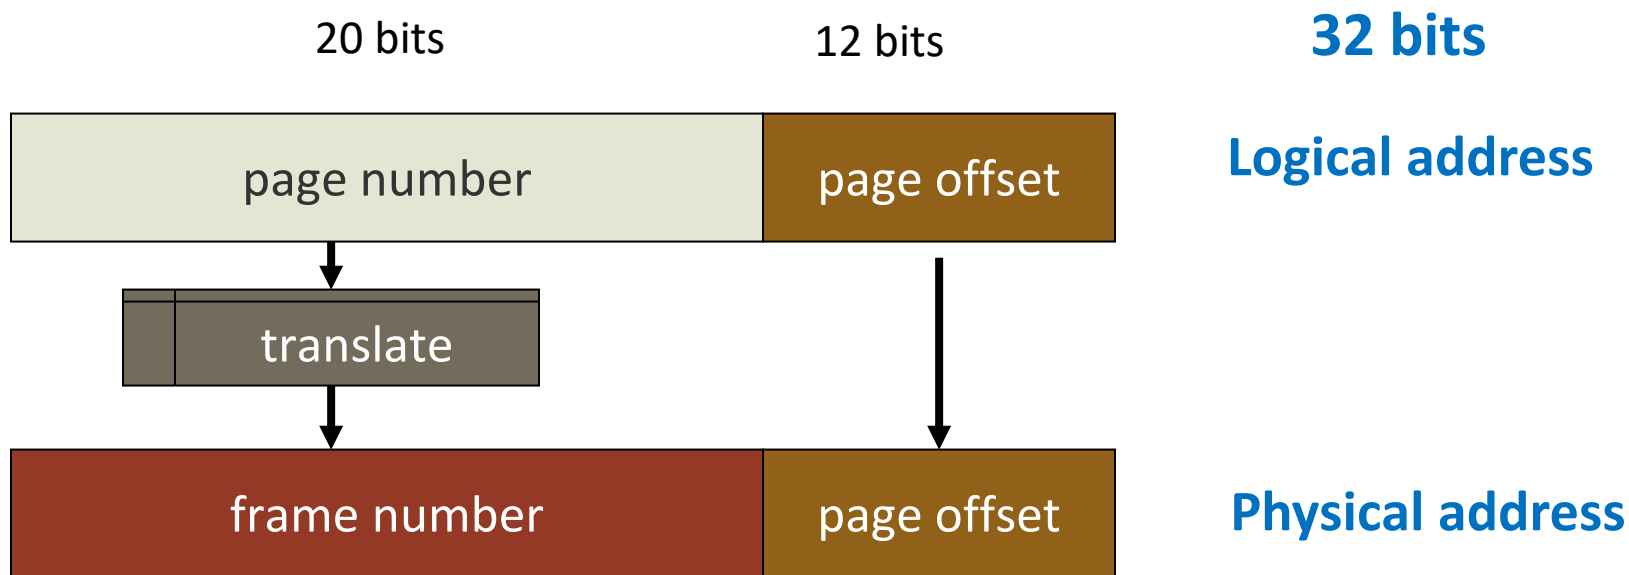
- **Goal:** **Eliminate** requirement that address space is **contiguous**
 - Eliminate external fragmentation
 - Grow segments as needed
- **Idea:** **Divide** address spaces and physical memory into **fixed-sized pages**
 - Size: 2^n , Example: 4KB
 - Physical page: page frame



Translation of Page Addresses

■ How to translate logical address to physical address?

- High-order bits of address designate page number
- Low-order bits of address designate offset within page




No addition needed; just append bits correctly...

How does format of address space determine number of pages and size of pages?


Virtual => Physical Page Mapping

Number of bits in
virtual address
format **does not**
need to equal
number of bits in
physical address
format



- How should OS translate VPN to PPN?
 - For segmentation, OS used a formula (e.g., $\text{phys addr} = \text{virt_offset} + \text{base_reg}$)
 - For paging, OS needs more general mapping mechanism
 - What data structure is good?
- Big array: pagetable

The Mapping



Where Are Pagetables Stored?

■ How big is a typical page table?

- assume 32-bit address space
- assume 4 KB pages
- assume 4 byte entries

■ Final answer: $2^{(32 - \log(4KB))} * 4 = 4 \text{ MB}$

- Page table size = Num entries * size of each entry
- Num entries = num virtual pages = $2^{(\text{bits for vpn})}$
- Bits for vpn = 32 – number of bits for page offset
 - $= 32 - \lg(4KB) = 32 - 12 = 20$
- Num entries = $2^{20} = 1 \text{ M}$
- Page table size = Num entries * 4 bytes = 4 MB

■ Implication: Store each page table in memory

- Hardware finds page table base with [register \(e.g., CR3 on x86\)](#)

■ What happens on a context-switch?

- [Change](#) contents of page table base register to newly scheduled process
- [Save](#) old page table base register in PCB of descheduled process

Other PT info

- **What other info is in pagetable entries besides translation?**
 - valid bit
 - protection bits
 - present bit (needed later)
 - reference bit (needed later)
 - dirty bit (needed later)
- **Pagetable entries are just bits stored in memory**
 - Agreement between hw and OS about interpretation

Memory Accesses with Pages

```
0x0010: movl (0x1100), %edi
0x0013: addl $0x3, %edi
0x0019: movl %edi, (0x1100)
```

Assume PT is at phys addr 0x5000

Assume PTE's are 4 bytes

Assume 4KB pages

How many bits for offset?

Simplified view
of page table

2
0
80
99

- **Old: How many mem refs with segmentation?**

5 (3 instrs, 2 movl)

Physical Memory Accesses with Paging?

1) Fetch instruction at logical addr 0x0010; vpn?

- Access page table to get ppn for vpn 0
- **Mem ref 1: 0x5000**
- **Learn vpn 0 is at ppn 2**
- Fetch instruction at 0x2010 (**Mem ref 2**)

Exec, load from logical addr 0x1100; vpn?

- Access page table to get ppn for vpn 1
- **Mem ref 3: 0x5004**
- Learn vpn 1 is at ppn 0
- **Movl from 0x0100 into reg (Mem ref 4)**

Pagetable is slow!!! Doubles memory references

Advantages of Paging

■ No external fragmentation

- Any page can be placed in any frame in physical memory

■ Fast to allocate and free

- Alloc: No searching for suitable free space
- Free: Doesn't have to coalesce with adjacent free space
- Just use bitmap to show free/allocated page frames

■ Simple to swap-out portions of memory to disk (later lecture)

- Page size matches disk block size
- Can run process when some pages are on disk
- Add “present” bit to PTE

Disadvantages of Paging

- **Internal fragmentation:** Page size may not match size needed by process
 - Wasted memory grows with larger pages
- **Additional memory reference to page table --> Very inefficient**
 - Page table must be stored in memory
 - MMU stores only base address of page table
 - **Solution: Add TLBs** (future lecture)
- **Storage for page tables may be substantial**
 - Simple page table: **Requires PTE for all pages** in address space
 - Entry needed even if page not allocated
 - Page tables must be **allocated contiguously** in memory
 - **Solution: Combine paging and segmentation** (future lecture)

