

Condition Variables

Questions answered in this lecture:

How can threads enforce ordering across operations?

How can `thread_join()` be implemented?

How can condition variables be used to support producer/consumer apps?

Concurrency Objectives

- **Mutual exclusion** (e.g., A and B don't run at the same time)
 - solved with **locks**
- **Ordering** (e.g., B runs after A does something)
 - solved with **condition variables** and **semaphores**

Ordering Example: Join

```
pthread_t p1, p2;  
Pthread_create(&p1, NULL, mythread, "A");  
Pthread_create(&p2, NULL, mythread, "B");  
// join waits for the threads to finish  
Pthread_join(p1, NULL);  
Pthread_join(p2, NULL);  
printf("main: done\n [balance: %d]\n [should: %d]\n",  
       balance, max*2);  
return 0;
```

how to implement join()?

Condition Variables

- **Condition Variable: queue of waiting threads**
- **B waits for a signal on CV before running**
 - `wait(CV, ...)`
- **A sends signal to CV when time for B to run**
 - `signal(CV, ...)`

Condition Variables

- **wait(cond_t *cv, mutex_t *lock)**
 - assumes the lock is held when wait() is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
 - This complexity stems from the desire to prevent certain race conditions from occurring when a thread is trying to put itself to sleep
- **signal(cond_t *cv)**
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return, doing nothing

Join Implementation: Attempt 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // x  
    Cond_wait(&c, &m);  // y  
    Mutex_unlock(&m);   // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);    // a  
    Cond_signal(&c);   // b  
    Mutex_unlock(&m);   // c  
}
```

Example schedule:

Parent: x y z

Child: a b c

Works!

Join Implementation: Attempt 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // x  
    Cond_wait(&c, &m);  // y  
    Mutex_unlock(&m);   // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);    // a  
    Cond_signal(&c);   // b  
    Mutex_unlock(&m);   // c  
}
```

Can you construct ordering that does not work?

Example broken schedule:

Parent: x y

Child: a b c

Parent waits forever!

Rule of Thumb 1

- **Keep state** in addition to CVs
- CVs are used to signal threads **when state changes**
- If state is already as needed, thread **doesn't wait** for a signal!

Join Implementation: Attempt 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // w  
    if (done == 0)    // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);  // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;    // a  
    Cond_signal(&c); // b  
}
```

Fixes previous broken ordering:

Parent: w x y z

Child: a b

Join Implementation: Attempt 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // w  
    if (done == 0)    // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);  // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;    // a  
    Cond_signal(&c); // b  
}
```

Can you construct ordering that does not work?

Parent: w x y

... sleep forever ...

Child: a b

Join Implementation: Correct

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // w  
    if (done == 0)    // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m); // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);    // a  
    done = 1;        // b  
    Cond_signal(&c);   // c  
    Mutex_unlock(&m); // d  
}
```

Parent: w x y z

Child: a b c

**hold the lock when calling signal
or wait, and you will always be
in good shape**


- Use mutex to ensure no race between interacting with state and wait/signal

Producer/Consumer Problem

Example: UNIX Pipes


- A pipe may have many **writers** and **readers**
- Internally, there is a **finite-sized buffer**
- **Writers add data to the buffer**
 - Writers have to wait if buffer is full
- **Readers remove data from the buffer**
 - Readers have to wait if buffer is empty

Example: UNIX Pipes




Example: UNIX Pipes

write!




Example: UNIX Pipes




Example: UNIX Pipes

write!




Example: UNIX Pipes




Example: UNIX Pipes

read!




Example: UNIX Pipes




Example: UNIX Pipes

write!




Example: UNIX Pipes




Example: UNIX Pipes

read!




Example: UNIX Pipes




Example: UNIX Pipes

read!




Example: UNIX Pipes




Example: UNIX Pipes

read!




Example: UNIX Pipes

read!




note: readers must wait

Example: UNIX Pipes




Example: UNIX Pipes

write!




Example: UNIX Pipes




Example: UNIX Pipes

write!



Example: UNIX Pipes



Example: UNIX Pipes

write!



Example: UNIX Pipes

write!

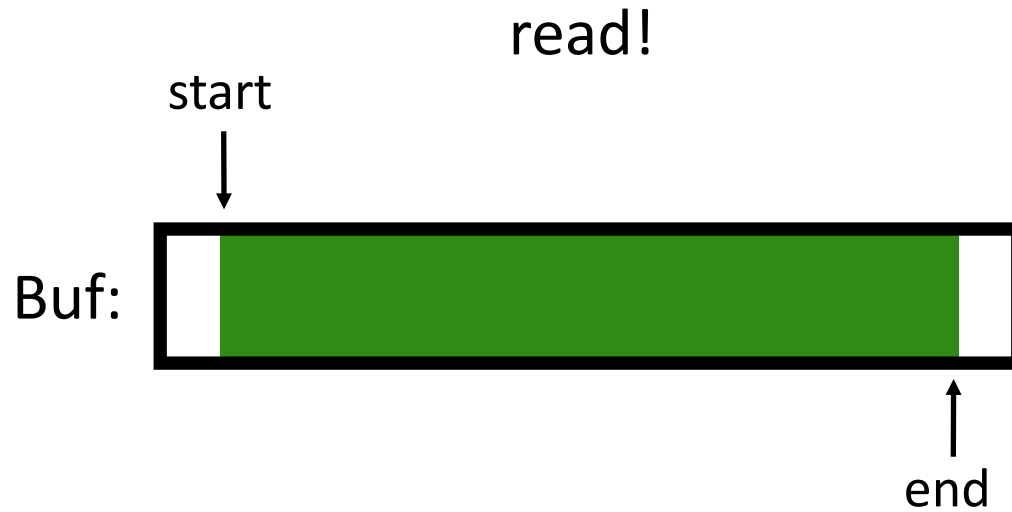


note: writers must wait

Example: UNIX Pipes



Example: UNIX Pipes



Example: UNIX Pipes

- **Implementation:**
 - reads/writes to buffer require locking
- **when buffers are full, writers must wait**
- **when buffers are empty, readers must wait**

Producer/Consumer Problem


- Producers generate data (like pipe writers)
- Consumers grab data and process it (like pipe readers)
- Producer/consumer problems are frequent in systems
- Web servers
 - General strategy use condition variables to:
 - make producers wait when buffers are full
 - make consumers wait when there is nothing to consume

Produce/Consumer Example

- **Start with easy case:**
 - 1 producer thread
 - 1 consumer thread
 - 1 shared buffer to fill/consume (max = 1)
- **Numfill = number of buffers currently filled**
- **Examine slightly broken code to begin...**


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

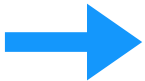
[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```



numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

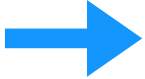
[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[SLEEPING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```



numfull=1

[SLEEPING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[SLEEPING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[SLEEPING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```



numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[RUNNING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNABLE]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=0

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


[SLEEPING]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```


numfull=1

[RUNNING]



```
void *producer(void *arg) {  
    for (int i=0; i<loops; i++) {  
        Mutex_lock(&m);  
        if (numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill(i);  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

[RUNNABLE]



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&cond, &m);  
        int tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

What about 2 consumers?

Can you find a problematic timeline with 2 consumers (still 1 producer)?

```

void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        Mutex_lock(&m); // p1
        if (numfull == 1) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(i); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}

```

```

void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        if (numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        printf("%d\n", tmp); // c7
    }
}

```

Producer:

Consumer1:

Consumer2:

c1 c2 c3

p1 p2 p4 p5 p6 p1 p2 p3
[Runnable]

another consumer
sneaks in

c1 c2 c4 c5 c6

get nothing!

c4


```

void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        Mutex_lock(&m); // p1
        while(numfull == 1) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(i); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}

```

```

void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        while(numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        printf("%d\n", tmp); // c7
    }
}

```

Customer1 wakes up and rechecks the state (c2)

```


void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        Mutex_lock(&m); // p1
        while(numfull == 1) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(i); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}

```

```

void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        while(numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        printf("%d\n", tmp); // c7
    }
}

```



Does last signal wake **producer** or **consumer1**?

How to wake the right thread?

- One solution:
wake all the threads!



Waking All Waiting Threads

- **wait(cond_t *cv, mutex_t *lock)**
 - assumes the lock is held when wait() is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- **signal(cond_t *cv)**
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return, doing nothing
- **broadcast(cond_t *cv)**
 - wake all waiting threads (if ≥ 1 thread is waiting)
 - if there are no waiting thread, just return, doing nothing

any disadvantage?

Example Need for Broadcast

```
void *allocate(int size) {  
    mutex_lock(&m);  
    while (bytesLeft < size)  
        cond_wait(&c);  
    ...  
}
```

```
void free(void *ptr, int size) {  
    ...  
    cond_broadcast(&c)  
    ...  
}
```

How to wake the right thread?

- **One solution:**

wake all the threads!



- **Better solution (usually): use two condition variables**

Producer/Consumer: Two CVs

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        Mutex_lock(&m); // p1  
        if (numfull == max) // p2  
            Cond_wait(&empty, &m); // p3  
        do_fill(i); // p4  
        Cond_signal(&fill); // p5  
        Mutex_unlock(&m); //p6  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if (numfull == 0)  
            Cond_wait(&fill, &m);  
        int tmp = do_get();  
        Cond_signal(&empty);  
        Mutex_unlock(&m);  
    }  
}
```

- Is this correct? Can you find a bad schedule?
 1. consumer1 waits because numfull == 0
 2. producer increments numfull, wakes consumer1
 3. before consumer1 runs, consumer2 runs, grabs entry, sets numfull=0.
 4. consumer1 then reads bad data.

Good Rule of Thumb 3

- Whenever a lock is acquired, **recheck assumptions about state!**
- Possible for another thread to grab lock in between signal and wakeup from wait
- Note that some libraries also have “spurious wakeups” (may wake multiple waiting threads at signal or at any time)

Producer/Consumer: Two CVs and While

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        Mutex_lock(&m); // p1  
        while (numfull == max) // p2  
            Cond_wait(&empty, &m); // p3  
        do_fill(i); // p4  
        Cond_signal(&fill); // p5  
        Mutex_unlock(&m); //p6  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        while (numfull == 0)  
            Cond_wait(&fill, &m);  
        int tmp = do_get();  
        Cond_signal(&empty);  
        Mutex_unlock(&m);  
    }  
}
```

- Is this correct? Can you find a bad schedule?
- **Correct!**
 - no concurrent access to shared state
 - every time lock is acquired, assumptions are reevaluated
 - a consumer will get to run after every do_fill()
 - a producer will get to run after every do_get()

Summary: rules of thumb for CVs

- **Keep state** in addition to CV's
- Always do wait/signal **with lock held**
- Whenever thread wakes from waiting, **recheck state**