

Semaphores

Questions answered in this lecture:

Review: How to implement join with condition variables?

Review: How to implement producer/consumer with condition variables?

What is the difference between **semaphores** and condition variables?

How to implement a **lock** with semaphores?

How to implement semaphores with locks and condition variables?

How to implement **join** and producer/consumer with semaphores?

How to implement **reader/writer locks** with semaphores?

Summary: rules of thumb for CVs

- Keep state in addition to CV's
- Always do wait/signal with lock held
- Whenever thread wakes from waiting, recheck state

Condition Variables vs Semaphores

- **Condition variables have no **state** (other than waiting queue)**
 - Programmer **must track additional state**
- **Semaphores have state: track integer value**
 - State cannot be directly accessed by user program, but state determines behavior of semaphore operations

Semaphore Operations

- **Allocate and Initialize**

```
sem_t sem;
```

```
sem_init(sem_t *s, int initval) {  
    s->value = initval;  
}
```

- User cannot read or write value directly after initialization

- **Wait or Test (sometime P() for Dutch word)**

- Waits until value of sem is > 0 , then **decrements** sem value

- **Signal or Increment or Post (sometime V() for Dutch)**

- **Increment** sem value, then wake a single waiter

wait and post are atomic

Join with CV vs Semaphores

CVs:

```
void thread_join() {  
    Mutex_lock(&m); // w  
    if (done == 0) // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m); // z  
}
```

```
void thread_exit() {  
    Mutex_lock(&m); // a  
    done = 1; // b  
    Cond_signal(&c); // c  
    Mutex_unlock(&m); // d  
}
```

Semaphores:

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

sem_t s;

sem_init(&s, ???);

Initialize to 0 (so sem_wait() must wait...)

```
void thread_join() {  
    sem_wait(&s);  
}
```

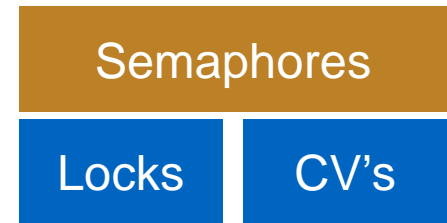
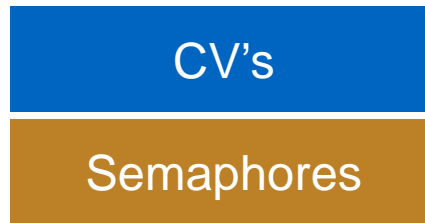
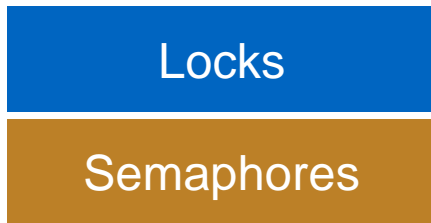
```
void thread_exit() {  
    sem_post(&s)  
}
```

Equivalence Claim

- Semaphores are equally powerful to Locks+CVs
 - what does this mean?
- One might be more convenient, but that's not relevant
- **Equivalence means each can be built from the other**

Proof Steps

- Want to show we can do these three things:



Build Lock from Semaphore

```
typedef struct __lock_t {  
    // whatever data structs you need go here  
} lock_t;
```

```
void init(lock_t *lock) {  
}
```

```
void acquire(lock_t *lock) {  
}
```

```
void release(lock_t *lock) {  
}
```

Sem_wait(): Waits until value > 0 , then decrement

Sem_post(): Increment value, then wake a single waiter

Locks

Semaphores

Build Lock from Semaphore

```
typedef struct __lock_t {  
    sem_t sem;  
} lock_t;  
  
void init(lock_t *lock) {  
    sem_init(&lock->sem, ??);  
}  
void acquire(lock_t *lock) {  
    sem_wait(&lock->sem);  
}  
void release(lock_t *lock) {  
    sem_post(&lock->sem);  
}
```

1 → 1 thread can grab lock

Sem_wait(): Waits until value > 0, then decrement

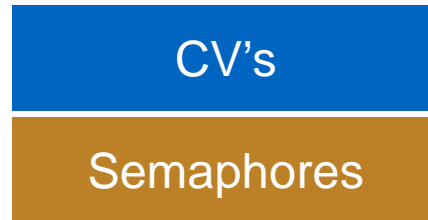
Sem_post(): Increment value, then wake a single waiter

Locks

Semaphores

Building CV's over Semaphores

- Possible, but really hard to do right



- Read about Microsoft Research's attempts:

<http://research.microsoft.com/pubs/64242/ImplementingCVs.pdf>

Build Semaphore from Lock and CV

```
Typedef struct {  
    // what goes here?
```

```
} sem_t;
```

```
Void sem_init(sem_t *s, int value) {  
    // what goes here?
```

```
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Typedef struct {  
    int value;  
    cond_t cond;  
    lock_t lock;  
} sem_t;  
  
Void sem_init(sem_t *s, int value) {  
    s->value = value;  
    cond_init(&s->cond);  
    lock_init(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    // what goes here?
```

```
Sem_post(sem_t *s) {  
    // what goes here?
```

```
}
```

```
}
```

Sem_wait(): Waits until value > 0 , then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    while (s->value <= 0)  
        cond_wait(&s->cond);  
    s->value--;  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    while (s->value <= 0)  
        cond_wait(&s->cond);  
    s->value--;  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    s->value++;  
    cond_signal(&s->cond);  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

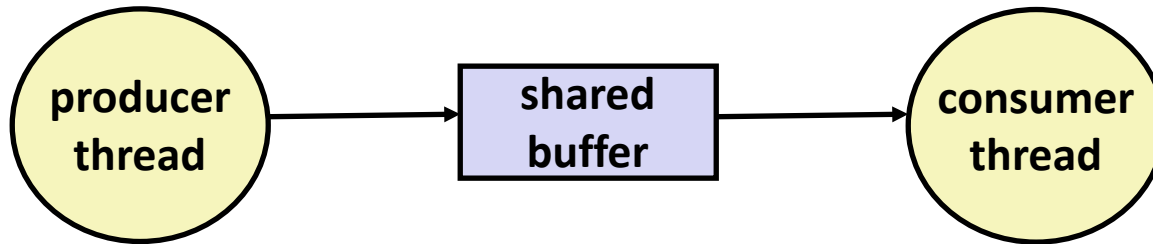
CV's

Using Semaphores to Coordinate Access to Shared Resources

- **Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true**
 - Use **counting semaphores** to keep track of resource state.
 - Use **binary semaphores** to notify other threads.
- **Two classic examples:**
 - The Producer-Consumer Problem
 - The Readers-Writers Problem



Producer-Consumer Problem



■ Common synchronization pattern:

- **Producer** waits for **empty slot**, inserts item in buffer, and notifies consumer
- **Consumer** waits for **item**, removes it from buffer, and notifies producer

■ Examples

- Multimedia processing:
 - Producer creates video frames, consumer renders them
- Event-driven graphical user interfaces
 - Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - Consumer retrieves events from buffer and paints the display

Producer/Consumer: Semaphores #1

■ Simplest case:

- Single producer thread, single consumer thread
- Single shared buffer with one element between producer and consumer

■ Requirements

- Consumer must wait for producer to fill buffer
- Producer must wait for consumer to empty buffer (if filled)

■ Requires 2 semaphores

- emptyBuffer: Initialize to ??? **1 → 1 empty buffer; producer can run 1 time first**
- fullBuffer: Initialize to ??? **0 → 0 full buffers; consumer can run 0 times first**

Producer

```
While (1) {  
  
    sem_wait(&emptyBuffer);  
    Fill(&buffer);  
  
    sem_signal(&fullBuffer);  
  
}
```

Consumer

```
While (1) {  
  
    sem_wait(&fullBuffer);  
    Use(&buffer);  
  
    sem_signal(&emptyBuffer);  
  
}
```

Producer/Consumer: Semaphores #2

■ Next case: Circular Buffer

- Single producer thread, single consumer thread
- Shared buffer with **N elements** between producer and consumer

■ Requires 2 semaphores

- emptyBuffer: Initialize to ??? **$N \rightarrow N$ empty buffers; producer can run N times first**
- fullBuffer: Initialize to ??? **$0 \rightarrow 0$ full buffers; consumer can run 0 times first**

Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    sem_signal(&fullBuffer);
}
```

Consumer

```
j = 0;
while (1) {
    sem_wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    sem_signal(&emptyBuffer);
}
```

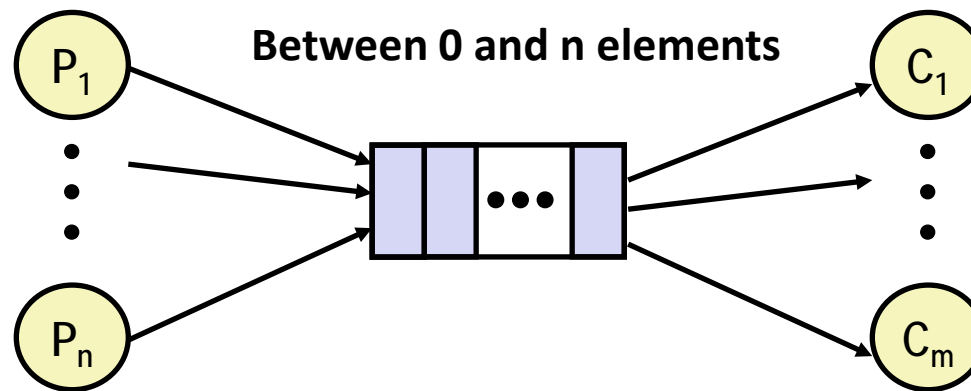
Producer/Consumer: Semaphore #3

■ Final case:

- Multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer

■ Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will previous code (shown below) not work???**



Producer/Consumer: Semaphore #3

■ Final case:

- Multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer

■ Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will previous code (shown below) not work???**

Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    sem_signal(&fullBuffer);
}
```

Consumer

```
j = 0;
while (1) {
    sem_wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    sem_signal(&emptyBuffer);
}
```

Are i and j private or shared? **Need each producer to grab unique buffer**

Producer/Consumer: Multiple Threads

■ Final case:

- Multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer

■ Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element

Producer

```
While (1) {  
    sem_wait(&emptyBuffer);  
    myi = findempty(&buffer);  
    Fill(&buffer[myi]);  
    sem_signal(&fullBuffer);  
}
```

Consumer

```
While (1) {  
    sem_wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    sem_signal(&emptyBuffer);  
}
```

Are myi and myj private or shared? Where is mutual exclusion needed???

Producer/Consumer: Multiple Threads

- Consider three possible locations for mutual exclusion
- Which work??? Which is best???

Producer #1

```
sem_wait(&mutex);  
sem_wait(&emptyBuffer);  
myi = findempty(&buffer);  
Fill(&buffer[myi]);  
sem_signal(&fullBuffer);  
sem_signal(&mutex);
```

Consumer #1

```
sem_wait(&mutex);  
sem_wait(&fullBuffer);  
myj = findfull(&buffer);  
Use(&buffer[myj]);  
sem_signal(&emptyBuffer);  
sem_signal(&mutex);
```

Problem?

Deadlock at mutex (e.g., consumer runs first; won't release mutex)

Producer/Consumer: Multiple Threads

- Consider three possible locations for mutual exclusion
- Which work??? Which is best???

Producer #2

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
Fill(&buffer[myi]);  
sem_signal(&mutex);  
sem_signal(&fullBuffer);
```

Consumer #2

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
Use(&buffer[myj]);  
sem_signal(&mutex);  
sem_signal(&emptyBuffer);
```

Works, but limits concurrency:

Only 1 thread at a time can be using or filling different buffers

Producer/Consumer: Multiple Threads

- Consider three possible locations for mutual exclusion
- Which work??? Which is best???

Producer #3

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
sem_signal(&mutex);  
Fill(&buffer[myi]);  
sem_signal(&fullBuffer);
```

Consumer #3

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
sem_signal(&mutex);  
Use(&buffer[myj]);  
sem_signal(&emptyBuffer);
```

Works and increases concurrency; only finding a buffer is protected by mutex;
Filling or Using different buffers can proceed concurrently

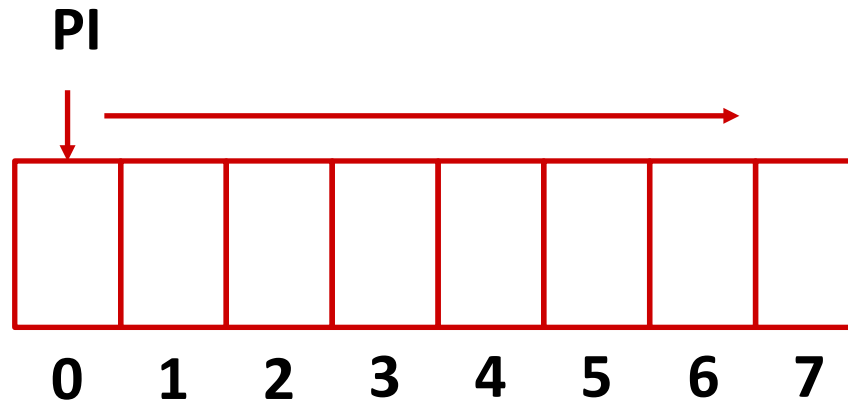
Producer/Consumer: Multiple Threads

- How to implement with multiple producer and multiple consumers?

Producer A:

`0 <- findempty`

An index PI
allocate
available slots



Producer/Consumer: Multiple Threads

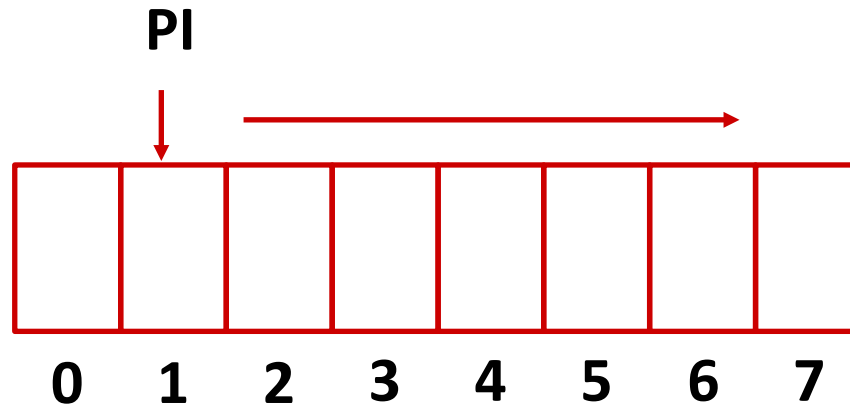
- How to implement with multiple producer and multiple consumers?

Producer A:

0 <- findempty

Producer B:

1 <- findempty

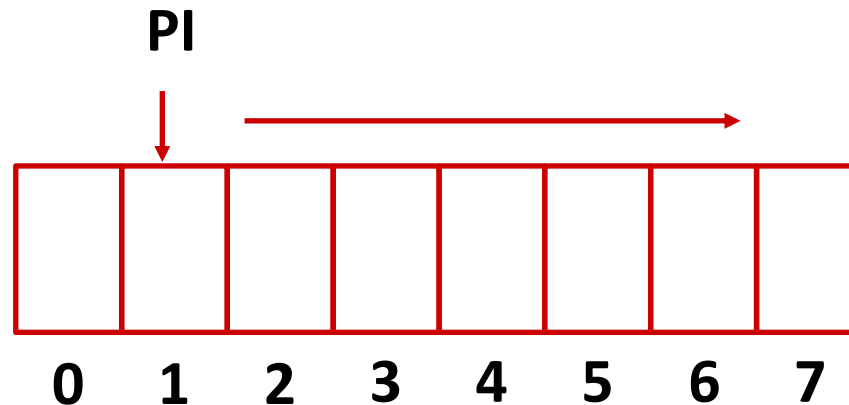


Producer/Consumer: Multiple Threads

- How to implement with multiple producer and multiple consumers?

Producer A:
0 <- findempty

Producer B:
1 <- findempty



Producer B
completes
signal(full_buffer)

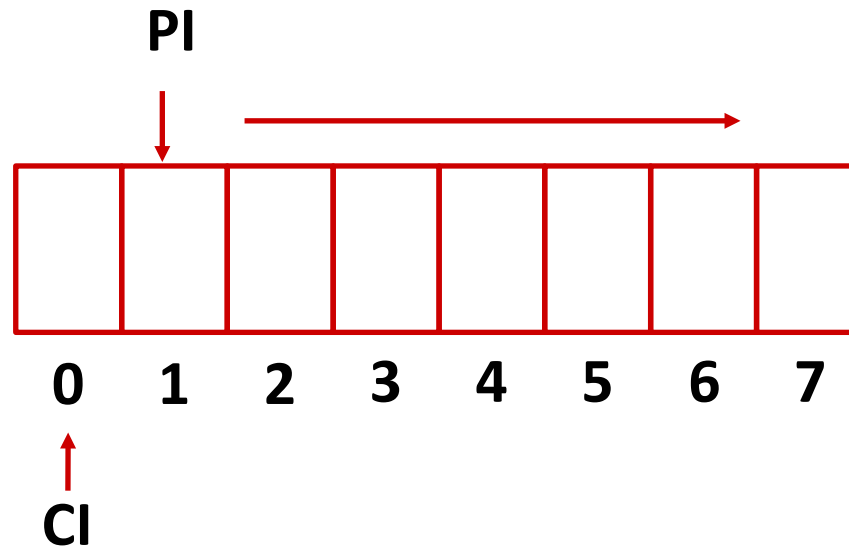
CI

which slot should a
consumer get?

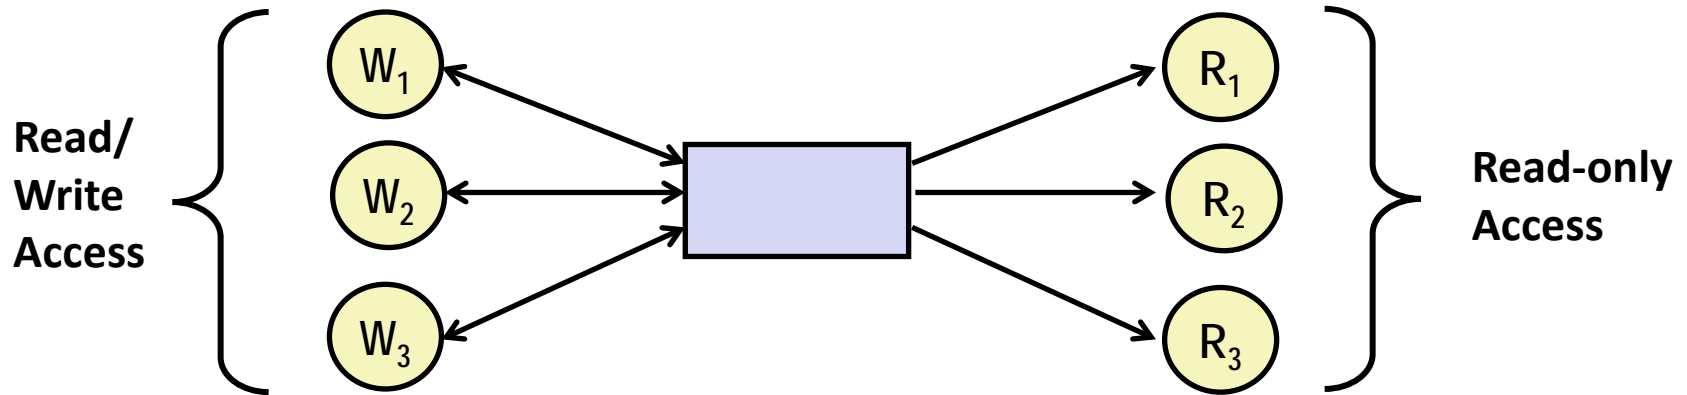
0 is incorrect
should return 1

Producer/Consumer: Multiple Threads

- How to implement with multiple producer and multiple consumers?
 - Approach 1: In find full, scan the buffer to find an available one
 - Approach 2: Producer insert an item into a list, which is searched by consumer



Readers-Writers Problem



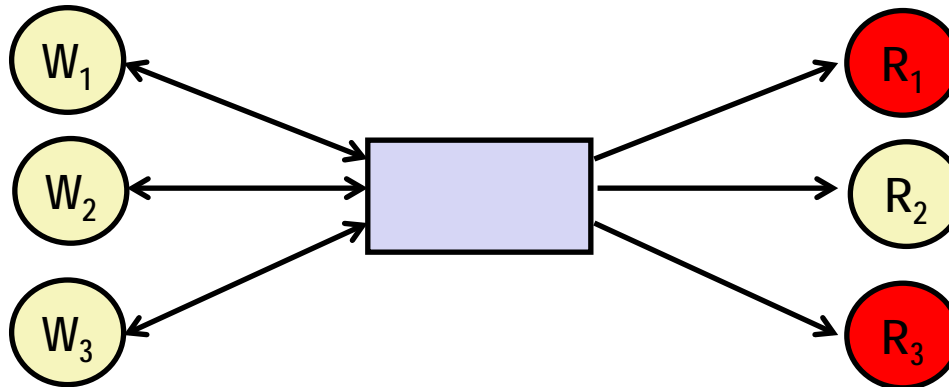
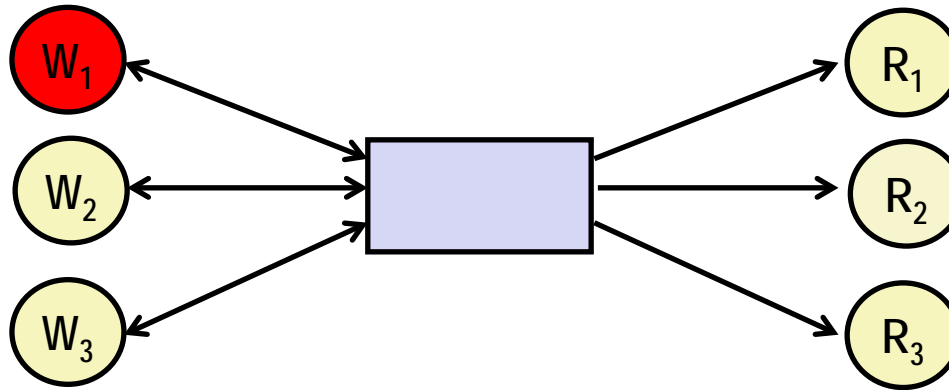
■ Problem statement:

- *Reader* threads only read the object
- *Writer* threads modify the object (read/write access)
- **Writers** must have **exclusive access** to the object
- **Unlimited** number of **readers** can access the object

■ Occurs frequently in real systems, e.g.,

- Online airline reservation system
- Multithreaded caching Web proxy

Readers/Writers Examples



Reader/Writer Locks

- Goal: Let **multiple reader** threads grab lock (**shared**)
- Only **one writer** thread can grab lock (**exclusive**)
 - No reader threads
 - No other writer threads
- How to implement reader/writer lock with semaphores
 - `rwlock_acquire_readlock`
 - `rwlock_release_readlock`
 - `rwlock_acquire_writelock`
 - `rwlock_release_writelock`

Variants of Readers-Writers

■ *First readers-writers problem* (favors readers)

- No reader should be kept waiting unless a writer has already been granted permission to use the object.
- A reader that arrives after a waiting writer gets priority over the writer.
- Web cache (写没那么重要)

■ *Second readers-writers problem* (favors writers)

- Once a writer is ready to write, it performs its write as soon as possible
- A reader that arrives after a writer must wait, even if the writer is also waiting.
- Ticket reservation

■ *Starvation* (where a thread waits indefinitely) is possible in both cases.

Reader/Writer Locks

```
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     //
15 }

21 void rwlock_release_readlock(rwlock_t *rw) {
22     //
23 }

29 rwlock_acquire_writelock(rwlock_t *rw) {
30     //
31 }

32 rwlock_release_writelock(rwlock_t *rw) {
33     //
34 }
```

T1: acquire_readlock()
T2: acquire_readlock()
T3: acquire_writelock()
T2: release_readlock()
T1: release_readlock()
T4: acquire_readlock()
T5: acquire_readlock() // ???
T3: release_writelock()
// what happens???

Reader/Writer Locks

```
1 typedef struct _rwlock_t {
2     sem_t lock;
3     sem_t writelock;
4     int readers;
5 } rwlock_t;
6
7 void rwlock_init(rwlock_t *rw) {
8     rw->readers = 0;
9     sem_init(&rw->lock, 1);
10    sem_init(&rw->writelock, 1);
11 }
12
```

Reader/Writer Locks

```
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock);
18     sem_post(&rw->lock);
19 }
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock);
26     sem_post(&rw->lock);
27 }
29 rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32 rwlock_release_writelock(rwlock_t *rw) {
33     sem_post(&rw->writelock);
34 }
```

T1: acquire_readlock()
T2: acquire_readlock()
T3: acquire_writelock()
T2: release_readlock()
T1: release_readlock()
T4: acquire_readlock()
T5: acquire_readlock() // ???
T3: release_writelock()
// what happens???

Solution to First Readers-Writers Problem

Readers: (读者永远 Block 写者)

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);
        // 第一个reader阻塞writer

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 1
W == 0

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        R2 → if (readcnt == 1) /* First in */
            P(&w);
            V(&mutex);

        R1 → /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 2
W == 0

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 2
W == 0

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 1
W == 0

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        R3 → if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        R2 → P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);

        R1 → }
    }
}
```

Writers:

```
void writer(void)
{
    while (1) { ← W1
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 2

W == 0

Solution to First Readers-Writers Problem

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w); ← W1

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 1
W == 0

Solution to First Readers-Writers Problem


Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */


        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```



R3

Writers:

```
void writer(void)
{
    while (1) {
        P(&w);
        /* Writing here */
        V(&w);
    }
}
```



W1

rw1.c

Arrivals: R1 R2 W1 R3

Readcnt == 0
W == 1

Other Versions of Readers-Writers

■ Shortcoming of first solution

- Continuous stream of readers will block writers indefinitely

■ Second version

- Once writer comes along, blocks access to later readers
- Series of writes could block all reads

■ FIFO implementation

- Service requests in order received
- Threads kept in FIFO
- Each has semaphore that enables its access to critical section

Solution to Second Readers-Writers Problem

- Consider two aspects
 - Following writes can overtake reads
 - Block multiple reads enter concurrently into the reading area

How to Modify the First Version to the Second

Readers:

```
int readcnt;      /* Initially 0 */
sem_t mutex, w; /* Both initially 1 */

void reader(void)
{
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&mutex);

        /* Reading happens here */

        P(&mutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&mutex);
    }
}
```

Writers:

```
void writer(void)
{
    while (1) {
        P(&w);

        /* Writing here */

        V(&w);
    }
}
```

rw1.c

Solution to Second Readers-Writers Problem

```
int readcnt, writecnt;           // Initially 0
sem_t rmutex, wmutex, r, w;     // Initially 1

void reader(void)
{
    while (1) {
        P(&r);
        P(&rmutex);
        readcnt++;
        if (readcnt == 1) /* First in */
            P(&w);
        V(&rmutex);
        V(&r);

        /* Reading happens here */

        P(&rmutex);
        readcnt--;
        if (readcnt == 0) /* Last out */
            V(&w);
        V(&rmutex);
    }
}
```

```
void writer(void)
{
    while (1) {
        P(&wmutex);
        writecnt++;
        if (writecnt == 1)
            P(&r);
        V(&wmutex);
        // 第一个writer阻塞reader

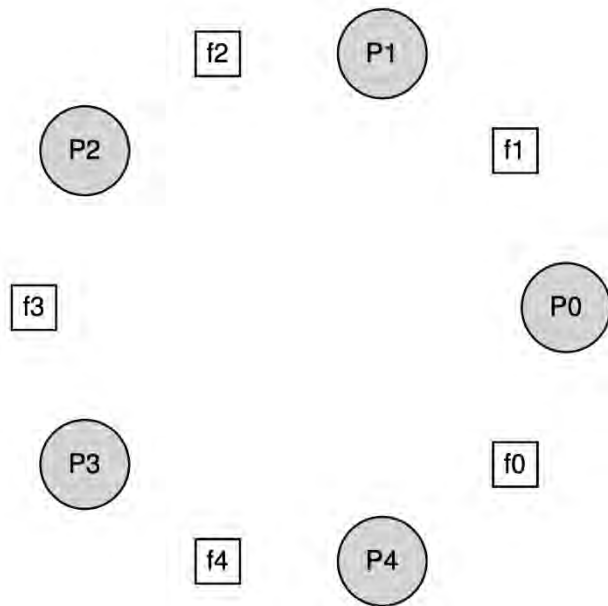
        P(&w);
        /* Writing here */
        V(&w);

        P(&wmutex);
        writecnt--;
        if (writecnt == 0);
            V(&r);
        V(&wmutex);
    }
}
```

Dining Philosopher's Problem

■ Problem Statement

- 5 philosophers on a table with 5 forks
- A philosopher needs 2 forks to eat
- Sometimes they think, no need for forks
- **Challenge: no deadlock and no philosopher starves**

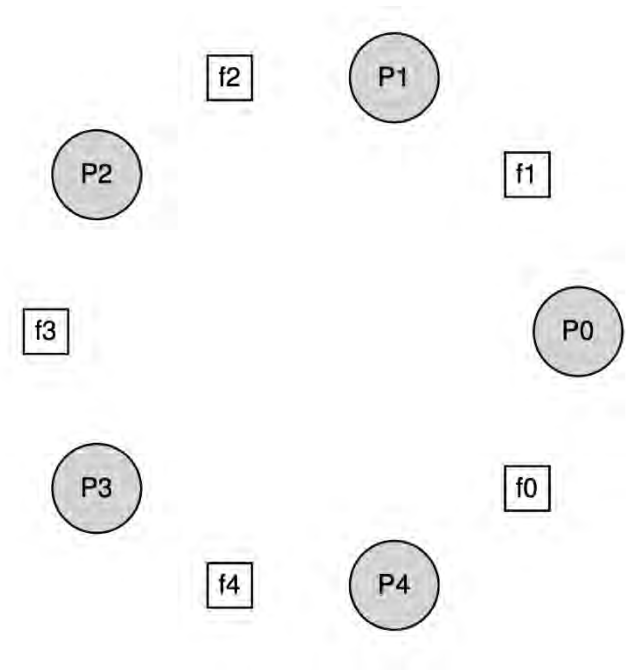


Each philosopher:

```
while (1) {  
    think();  
    get_forks(p);  
    eat();  
    put_forks(p);  
}
```

First Try

```
void get_forks(int p) {  
    sem_wait(&forks[left(p)]);  
    sem_wait(&forks[right(p)]);  
}  
  
void put_forks(int p) {  
    sem_post(&forks[left(p)]);  
    sem_post(&forks[right(p)]);  
}
```

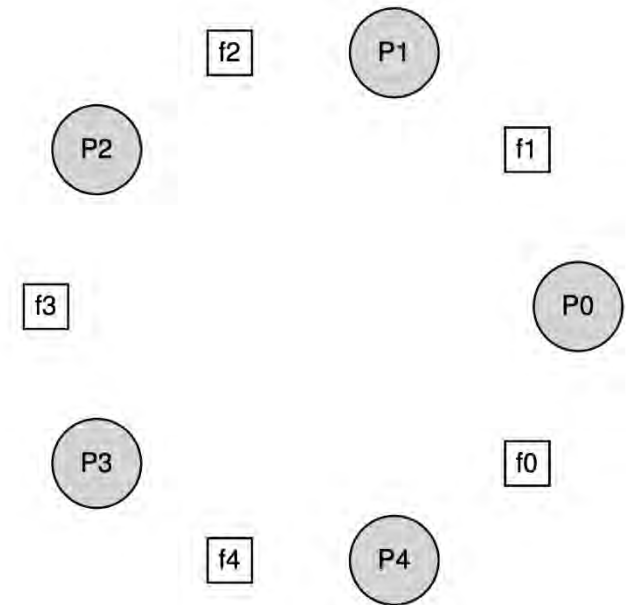


■ Wrong and cause deadlock

- If every philosopher takes their left fork simultaneously, no one can get the right fork
- Deadlock!

Second Try: break the dependency

```
void get_forks(int p) {  
    if (p == 4) {  
        sem_wait(&forks[right(p)]);  
        sem_wait(&forks[left(p)]);  
    } else {  
        sem_wait(&forks[left(p)]);  
        sem_wait(&forks[right(p)]);  
    }  
}
```



- **Cycle of waiting is broken!**
 - More deadlocks on next lecture

Semaphores

- **Semaphores are equivalent to locks + condition variables**
 - Can be used for both mutual exclusion and ordering
- **Semaphores contain state**
 - How they are initialized depends on how they will be used
 - Init to 1: Mutex
 - Init to 0: Join (1 thread must arrive first, then other)
 - Init to N: Number of available resources
- **Sem_wait(): Waits until value > 0, then decrement (atomic)**
- **Sem_post(): Increment value, then wake a single waiter (atomic)**
- **Can use semaphores in producer/consumer relationships and for reader/writer locks**