

Concurrency Bugs

Questions answered in this lecture:

Why is concurrent programming difficult?

What type of concurrency bugs occur?

How to fix **atomicity bugs** (with locks)?

How to fix **ordering bugs** (with condition variables)?

How does **deadlock** occur?

How to prevent deadlock (with waitfree algorithms, grab all locks atomically, trylocks, and ordering across locks)?

Concurrency in Medicine: Therac-25 (1980's)

- The **Therac-25** was a computer-controlled [radiation therapy](#) machine produced by [Atomic Energy of Canada Limited](#) (AECL) in 1982
- It was involved in at least six accidents between 1985 and 1987, in which patients were given massive [overdoses of radiation](#).
- Because of [concurrent programming errors](#) (also known as race conditions), it sometimes gave its patients radiation doses that were hundreds of times greater than normal, resulting in death or serious injury.
- Become a standard case study in [health informatics](#), [software engineering](#), and [computer ethics](#).

Source: <http://en.wikipedia.org/wiki/Therac-25>


Concurrency in Medicine: Therac-25 (1980's)

“The accidents occurred when the **high-power electron beam** was activated instead of the intended low power beam, and without the beam spreader plate rotated into place. Previous models had **hardware interlocks** in place to prevent this, but Therac-25 had **removed** them, depending instead on software interlocks for safety. The software interlock could fail due to a **race condition**.”

“...in three cases, the injured patients **later died**.”

Source: <http://en.wikipedia.org/wiki/Therac-25>

Concurrency Study from 2008



Lu *etal.* Study:

For four major projects, search for **concurrency bugs** among >500K bug reports. Analyze small sample to identify common types of concurrency bugs.

Source: <http://pages.cs.wisc.edu/~shanlu/paper/asplos122-lu.pdf>

Atomicity: MySQL

Thread 1:

```
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info, ...);  
    ...  
}
```

Thread 2:

```
thd->proc_info = NULL;
```

What's wrong?

Test (`thd->proc_info != NULL`) and set (writing to `thd->proc_info`) should be atomic

Fix Atomicity Bugs with Locks


Thread 1:

```
pthread_mutex_lock(&lock);  
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info, ...);  
    ...  
}  
pthread_mutex_unlock(&lock);
```

Thread 2:

```
pthread_mutex_lock(&lock);  
thd->proc_info = NULL;  
pthread_mutex_unlock(&lock);
```

Concurrency Study from 2008



Lu *etal.* Study:

For four major projects, search for **concurrency bugs** among >500K bug reports. Analyze small sample to identify common types of concurrency bugs.

Source: <http://pages.cs.wisc.edu/~shanlu/paper/asplos122-lu.pdf>

Ordering: Mozilla

Thread 1:

```
void init() {  
    ...  
    mThread =  
        PR_CreateThread(mMain, ...);  
    ...  
}
```

Thread 2:

```
void mMain(...) {  
    ...  
    mState = mThread->State;  
    ...  
}
```

What's wrong?

Thread 1 sets value of mThread needed by Thread2

How to ensure that reading mThread happens after mThread initialization?

Fix Ordering bugs with Condition variables

```
Thread 1:
void init() {
    ...

    mThread =
        PR_CreateThread(mMain, ...);

    Mutex_lock(&mtLock);
    mtInit = 1;
    pthread_cond_signal(&mtCond);
    Mutex_unlock(&mtLock);

    ...
}
```

```
Thread 2:

void mMain(...) {
    ...

    Mutex_lock(&mtLock);
    while (mtInit == 0)
        Cond_wait(&mtCond, &mtLock);
    Mutex_unlock(&mtLock);

    mState = mThread->State;
    ...
}
```



One Worry: Races

- A *race* occurs when correctness of the program depends on one thread reaching point x before another thread reaches point y

```
/* a threaded program with a race */
int main(int argc, char** argv) {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++)
        Pthread_create(&tid[i], NULL, thread, &i);
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    return 0;
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```




One Worry: Races

- A *race* occurs when correctness of the program depends on one thread reaching point x before another thread reaches point y

```
/* a threaded program with a race */
int main(int argc, char** argv) {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++)
        Pthread_create(&tid[i], NULL, thread, &i);
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    return 0;
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

Data Race



Race Elimination

- Make sure **don't** have **unintended sharing** of state

```
/* a threaded program without the race */
int main(int argc, char** argv) {
    pthread_t tid[N];
    int i;
    for (i = 0; i < N; i++) {
        int *valp = Malloc(sizeof(int));
        *valp = i;
        Pthread_create(&tid[i], NULL, thread, valp);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    return 0;
}

/* thread routine */
void *thread(void *vargp) {
    int myid = *((int *)vargp);
    Free(vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

[norace.c](#)


Race Elimination

■ Is this correct?

```
/* a threaded program without the race */
int main(int argc, char** argv) {
    pthread_t tid[N];
    uint64_t i;
    for (i = 0; i < N; i++) {
        Pthread_create(&tid[i], NULL, thread, (void *)i);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tid[i], NULL);
    return 0;
}

/* thread routine */
void *thread(void *vargp) {
    uint64_t myid = (uint64_t)vargp;
    Free(vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

Concurrency Study from 2008



Lu *etal.* Study:

For four major projects, search for **concurrency bugs** among >500K bug reports. Analyze small sample to identify common types of concurrency bugs.

Source: <http://pages.cs.wisc.edu/~shanlu/paper/asplos122-lu.pdf>


Deadlock


Deadlock: **No progress can be made** because two or more threads are waiting for the other to take some action and thus neither ever does


“Cooler” name: the **deadly embrace** (Dijkstra)


Deadlock


- Another Def: A process is *deadlocked* if it is waiting for a condition that will never be true.
- Typical Scenario
 - Processes 1 and 2 needs two resources (A and B) to proceed
 - Process 1 acquires A, waits for B
 - Process 2 acquires B, waits for A
 - Both will wait forever!
















who goes?














who goes?



who goes?





Code Example

Thread 1:


```
lock(&A);  
lock(&B);
```

Thread 2:


```
lock(&B);  
lock(&A);
```

Can deadlock happen with these two threads?

Circular Dependency



Deadlock



Fix Deadlocked Code

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

How would you fix this code?


Thread 1

```
lock(&A);  
lock(&B);
```

Thread 2

```
lock(&A);  
lock(&B);
```

Non-circular Dependency (fine)



What's Wrong?

```
set_t *set_intersection (set_t *s1, set_t *s2) {  
    set_t *rv = Malloc(sizeof(*rv));  
    Mutex_lock(&s1->lock);  
    Mutex_lock(&s2->lock);  
    for(int i=0; i<s1->len; i++) {  
        if(set_contains(s2, s1->items[i])  
            set_add(rv, s1->items[i]);  
    }  
    Mutex_unlock(&s2->lock);  
    Mutex_unlock(&s1->lock);  
}
```

Encapsulation

- Modularity can make it **harder to see deadlocks**

Thread 1:

```
rv = set_intersection(setA,  
                      setB);
```

Thread 2:

```
rv = set_intersection(setB,  
                      setA);
```

Solution?

```
if (m1 > m2) { // Compare some kind of encoding  
    // grab locks in high-to-low address order  
    pthread_mutex_lock(m1);  
    pthread_mutex_lock(m2);  
} else {  
    pthread_mutex_lock(m2);  
    pthread_mutex_lock(m1);  
}
```


Any other problems?

Code assumes $m1 \neq m2$ (not same lock)

Deadlock Theory

- **Deadlocks can only happen with these four conditions:**

- mutual exclusion
- hold-and-wait
- no preemption
- circular wait



- **Eliminate deadlock by eliminating any one condition**

Mutual Exclusion

- **Def:**
- **Threads claim exclusive control of resources that they require (e.g., thread grabs a lock)**

Wait-Free Algorithms

- Strategy: **Eliminate locks!**
- Try to replace locks with **atomic primitive**:
- **int CompAndSwap(int *addr, int expected, int new)**
Returns 0: fail, 1: success

```
void add (int *val, int amt) {  
    Mutex_lock(&m);  
    *val += amt;  
    Mutex_unlock(&m);  
}
```

```
void add (int *val, int amt) {  
    do {  
        int old = *value;  
    } while(!CompAndSwap(val, ??, old+amt));  
}
```

?? → old

Wait-Free Algorithms: Linked List Insert

- Strategy: **Eliminate locks!**
- Try to replace locks with **atomic primitive**:
- **int CompAndSwap(int *addr, int expected, int new)**
Returns 0: fail, 1: success

```
void insert (int val) {  
    node_t *n = Malloc(sizeof(*n));  
    n->val = val;  
    lock(&m);  
    n->next = head;  
    head = n;  
    unlock(&m);  
}
```

```
void insert (int val) {  
    node_t *n = Malloc(sizeof(*n));  
    n->val = val;  
    do {  
        n->next = head;  
    } while (!CompAndSwap(&head,  
        n->next, n));  
}
```


Deadlock Theory

- **Deadlocks can only happen with these four conditions:**
 - mutual exclusion
 - hold-and-wait
 - no preemption
 - circular wait
- **Eliminate deadlock by eliminating any one condition**

Hold-and-Wait

- **Def:**
- **Threads hold resources allocated to them (e.g., locks they have already acquired) while waiting for additional resources (e.g., locks they wish to acquire).**

Eliminate Hold-and-Wait

- **Strategy: Acquire all locks atomically once**

- Can release locks over time, but cannot acquire again until all have been released
- How to do this?

- **Use a meta lock, like this:**

```
lock(&meta);
```

```
lock(&L1);
```

```
lock(&L2);
```

```
...
```

```
unlock(&meta);
```

```
// Critical section code
```

```
unlock(...);
```

Disadvantages?

Must know ahead of time which locks will be needed

Must be conservative (acquire any lock possibly needed)

Degenerates to just having one big lock

Deadlock Theory

- **Deadlocks can only happen with these four conditions:**
 - mutual exclusion
 - hold-and-wait
 - no preemption
 - circular wait
- **Eliminate deadlock by eliminating any one condition**

No preemption

- Def:
- Resources (e.g., locks) cannot be forcibly removed from threads that are holding them.

Support Preemption

- Strategy: if thread can't get what it wants, release what it holds

top:

```
lock(A);  
if (trylock(B) == -1) {  
    unlock(A);  
    goto top;  
}  
...
```

Disadvantages?

- Livelock: no processes make progress, but the state of involved processes constantly changes
How to address?

Support Preemption

- Classic solution: **Exponential back-off** -- wait some time

```
    i = 1;
top:
    lock(A);
    if (trylock(B) == -1) {
        unlock(A);
        sleep(i);
        i *= 2;
        goto top;
    }
    ...
```

Deadlock Theory

- **Deadlocks can only happen with these four conditions:**
 - mutual exclusion
 - hold-and-wait
 - no preemption
 - circular wait
- **Eliminate deadlock by eliminating any one condition**

Circular Wait

- Def:
- There exists a **circular chain of threads** such that each thread holds a resource (e.g., lock) being requested by next thread in the chain.

Eliminating Circular Wait

■ Strategy:

- decide which **locks should be acquired before others**
- if A before B, never acquire A if B is already held!
- **document** this, and write code accordingly

■ Works well if system has distinct layers

Releasing Lock in Other Order

thread 0

```
A.lock();
```

```
B.lock();
```

```
B.unlock();
```

```
A.unlock();
```

thread 1

```
A.lock();
```

```
B.lock();
```

```
A.unlock();
```

```
B.unlock();
```

Releasing Lock in Other Order

thread 0

```
A.lock();  
B.lock();  
Foo();  
B.unlock();  
Bar();  
A.unlock();
```

thread 1

```
A.lock();  
B.lock();  
Foo();  
A.unlock();  
Bar();  
B.unlock();
```

Any problem?

Releasing Lock in Other Order

thread 0

```
A.lock();  
B.lock();  
Foo();  
B.unlock();  
Bar();  
A.unlock();
```

thread 1

```
A.lock();  
B.lock();  
Foo();  
A.unlock();  
Bar();  
B.unlock();
```

If **Bar()** attempts to reacquire A, you've effectively broken your lock ordering. You're holding B and then trying to get A. Now it **can** deadlock.

Lock Ordering in Linux

In linux-3.2.51/include/linux/fs.h

```
/* inode->i_mutex nesting subclasses for the lock
```

```
* validator:
```

```
* 0: the object of the current VFS operation
```

```
* 1: parent
```

```
* 2: child/target
```

```
* 3: quota file
```

```
* The locking order between these classes is
```

```
* parent -> child -> normal -> xattr -> quota
```

```
*/
```

Banker's Algorithm

- Banker's algorithm is a resource allocation and deadlock avoidance algorithm developed by Edsger Dijkstra that tests for safety by simulating the allocation of predetermined **maximum possible amounts of all resources**, and then makes an "s-state" check to test for possible deadlock conditions for all other pending activities, before deciding whether allocation should be allowed to continue.
- When a new process enters a system, it must declare the **maximum number of instances of each resource** type that it may ever claim; clearly, that number may **not exceed the total** number of resources in the system. Also, when a process gets all its requested resources it must **return them in a finite amount of time**.

Banker's Algorithm

- For the Banker's algorithm to work, it needs to know three things:
 - How much of each resource each process could possibly request ("MAX")
 - How much of each resource each process is currently holding ("ALLOCATED")
 - How much of each resource the system currently has available ("AVAILABLE")
- Resources (memory, semaphores, interface) may be allocated to a process only if the amount of resources requested is less than or equal to the amount available; otherwise, the process waits until resources are available.

Banker's Algorithm

- The Banker's algorithm derives its name from the fact that this algorithm could be used in a **banking system** to ensure that the bank does not **run out of resources**, because the bank would never allocate its money in such a way that it can no longer satisfy the needs of all its customers.
- By using the Banker's algorithm, the bank ensures that when customers request money **the bank never leaves a safe state**. If the customer's request does not cause the bank to leave a safe state, the cash will be allocated, otherwise the customer **must wait until some other customer deposits enough**.

Banker's Algorithm

Total system resources:

A	B	C	D
6	5	7	6

Available system res:

A	B	C	D
3	1	1	2

Needed resources:

max - allocated

	A	B	C	D
P1	2	1	0	1
P2	0	2	0	1
P3	0	1	4	0

Current allocated

resources for processes:

	A	B	C	D
P1	1	2	2	1
P2	1	0	3	3
P3	1	2	1	0

Maximum allocated

resources for processes:

	A	B	C	D
P1	3	3	2	2
P2	1	2	3	4
P3	1	3	5	0

Banker's Algorithm – Safe State

Available system res:

A	B	C	D
3	1	1	2

Current allocated
resources for processes:

	A	B	C	D
P1	1	2	2	1
P2	1	0	3	3
P3	1	2	1	0

Needed resources:
(max – allocated)

	A	B	C	D
P1	2	1	0	1
P2	0	2	0	1
P3	0	1	4	0

Assumption: the system assumes that all processes will eventually attempt to acquire their stated maximum resources and terminate soon afterward.

Safe state: a state is considered safe if it is possible for all processes to finish executing (terminate).

Safe state?

Banker's Algorithm – P1-P2-P3

Available system res:

A	B	C	D
3	1	1	2

Current allocated
resources for processes:

	A	B	C	D
P1	1	2	2	1
P2	1	0	3	3
P3	1	2	1	0

Needed resources:
(max – allocated)

	A	B	C	D
P1	2	1	0	1
P2	0	2	0	1
P3	0	1	4	0

Available system res after P1:

A	B	C	D
4	3	3	3

After P1
resources for processes:

	A	B	C	D
P1				
P2	1	0	3	3
P3	1	2	1	0

Safe!

Banker's Algorithm – P1-P2-P3

Available system res after P1:

A	B	C	D
4	3	3	3

Available system res after P2:

A	B	C	D
5	3	6	6

After P1

resources for processes:

	A	B	C	D
P1				
P2	1	0	3	3
P3	1	2	1	0

After P2

resources for processes:

	A	B	C	D
P1				
P2				
P3	1	2	1	0

Needed resources:

(max – allocated)

	A	B	C	D
P1	2	1	0	1
P2	0	2	0	1
P3	0	1	4	0

Safe!

Banker's Algorithm – Unsafe State

Available system res:

A	B	C	D
3	1	1	2

Current allocated
resources for processes:

	A	B	C	D
P1	1	2	2	1
P2	1	0	3	3
P3	1	2	1	0

Needed resources:
(max – allocated)

	A	B	C	D
P1	2	1	0	1
P2	0	2	0	1
P3	0	1	4	0

Assumption: the system assumes that all processes will eventually attempt to acquire their stated maximum resources and terminate soon afterward.

Safe state: a state is considered safe if it is possible for all processes to finish executing (terminate).

Unsafe state?

Summary

- When in doubt about correctness, better to limit concurrency (i.e., **add unnecessary lock**)
- Concurrency is hard, encapsulation makes it harder!
- Have a strategy to avoid deadlock and stick to it
- Choosing a **lock order** is probably most practical