

File System Implementation

Questions answered in this lecture:

What on-disk structures to represent files and directories?

Contiguous, Extents, Linked, FAT, Indexed, Multi-level indexed

Which are good for different metrics?

What disk operations are needed for:

- make directory

- open file

- write/read file

- close file

Review: File Names

- **Different types of names work better in different contexts**
- **inode**
 - unique name for file system to use
 - records meta-data about file: file size, permissions, etc
- **path**
 - easy for people to remember
 - organizes files in hierarchical manner; encode locality information
- **file descriptor**
 - avoid frequent traversal of paths
 - remember multiple offsets for next read or write

Review: File API

`int fd = open(char *path, int flag, mode_t mode)`

`read(int fd, void *buf, size_t nbyte)`

`write(int fd, void *buf, size_t nbyte)`

`close(int fd)`

Implementation

■ 1. On-disk structures

- how does file system represent files, directories?



■ 2. Access methods

- what steps must reads/writes take?

Part 1: Disk Structures

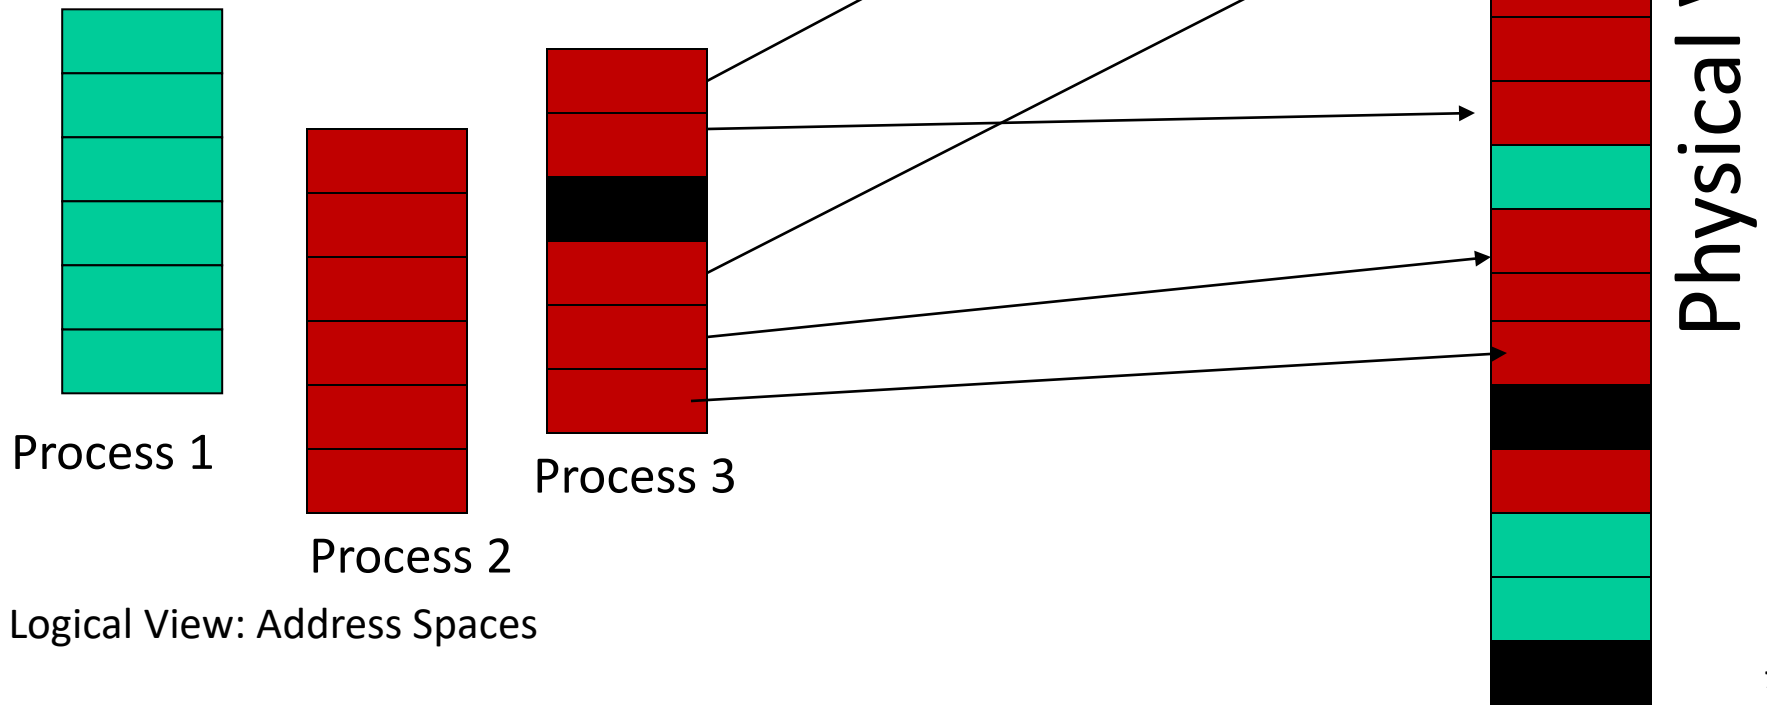
Persistent Store

- **Given:** large array of blocks on disk
- **Want:** some structure to map files to disk blocks



Similarity to Memory?

**Same principle:
map logical abstraction to physical resource**



Allocation Strategies

■ Many different approaches

- Contiguous
- Extent-based
- Linked
- File-allocation Tables
- Indexed
- Multi-level Indexed


■ Questions

- Amount of **fragmentation** (internal and external)
 - freespace that can't be used
- Ability to **grow** file over time?
- Performance of **sequential accesses** (contiguous layout)?
- Speed to find data blocks for **random accesses**?
- Wasted space for **meta-data overhead** (everything that isn't data)?
 - Meta-data must be stored persistently too!

Contiguous Allocation

■ Allocate each file to contiguous sectors on disk

- Meta-data: Starting block and size of file
- OS allocates by finding sufficient free space
 - Must predict future size of file; Should space be reserved?
- Example: IBM OS/360



Fragmentation (internal and external)?

- Horrible **external fragmentation** (needs periodic compaction)

Ability to grow file over time?

- **May not be able to without moving**

Seek cost for sequential accesses?

+ Excellent performance

Speed to calculate random accesses?

+ Simple calculation

Wasted space for meta-data?


+ Little overhead for meta-data



Small Fixed Number of ExtentsS

■ Allocate **multiple contiguous regions (extents)** per file

- Meta-data: Small array (2-6) designating each extent
Each entry: starting block and size



Fragmentation (internal and external)?

- Helps external fragmentation

Ability to grow file over time?

- Can grow (until run out of extents)

Seek cost for sequential accesses?

+ Still good performance

Speed to calculate random accesses?

+ Still simple calculation

Wasted space for meta-data?


+ Still small overhead for meta-data



Linked Allocation

■ Allocate **linked-list of fixed-sized blocks** (multiple sectors)

- Meta-data: Location of first block of file
Each block also contains pointer to next block
- Examples: TOPS-10, Alto



Fragmentation (internal and external)?

+ No external frag (use any block); internal: pointer and space in last block

Ability to grow file over time?

+ Can grow easily

Seek cost for sequential accesses?

+/- Depends on data layout

Speed to calculate random accesses?

- Ridiculously poor

Wasted space for meta-data?


- Waste pointer per block

Trade-off: Block size (does not need to equal sector size)

File-Allocation Table (FAT)

■ Variation of Linked allocation

- Keep linked-list information for all files in on-disk FAT table
- Meta-data: Location of first block of file
 - And, FAT table itself



- Draw corresponding FAT Table?
- Comparison to Linked Allocation
 - Same basic advantages and disadvantages
 - Disadvantage: Read from two disk locations for every data read
 - Optimization:
 - Cache FAT in main memory
 - Advantage: Greatly improves random accesses
 - What portions should be cached? Scale with larger file systems?

File-Allocation Table (FAT)

File Allocation Table

File.txt



is

Block 2




Block 6



Block 3



Block 5



	<u>FAT</u>	
	Busy	Next
0	0	
1	1	-1
2	1	6
3	1	5
4	1	-1
5	1	-1
6	1	3
7	0	

Directory Table Format

filename	starting block	meta data
foo	1	

/foo

filename	starting block	meta data
File.txt	2	

Indexed Allocation

- **Allocate fixed-sized blocks for each file**
 - Meta-data: **Fixed-sized array of block pointers**
 - Allocate space for ptrs at file creation time



- **Advantages**
 - No external fragmentation
 - Files can be easily grown up to max file size
 - Supports random access
- **Disadvantages**
 - **Large overhead for meta-data:**
 - Wastes space for unneeded pointers (most files are small!)


Inode


type
uid
rwx
size
blocks
time
ctime
links_count
addrs[N]

- Assume **single level** (just pointers to data blocks)
- What is **max file size**?
 - Assume **256-byte** inodes (all can be used for pointers)
 - Assume **4-byte** addrs
- How to get larger files?

$$256 / 4 = 64$$

$$64 * 4K = 256 \text{ KB!}$$





Indirect blocks are stored in regular data blocks.

what if we want to
[optimize for small files?](#)


Multi-Level Indexing

■ Optimize for small files

- Most files are small
- Based on workloads - an important design principle

Most files are small	~2K is the most common size
Average file size is growing	Almost 200K is the average
Most bytes are stored in large files	A few big files use most of space
File systems contains lots of files	Almost 100K on average
File systems are roughly half full	Even as disks grow, file systems remain ~50% full
Directories are typically small	Many have few entries; most have 20 or fewer

Figure 40.2: File System Measurement Summary




Better for small files


Multi-Level Indexing

■ Variation of Indexed Allocation

- Dynamically allocate **hierarchy of pointers** to blocks as needed
- Meta-data: Small number of pointers allocated statically
 - Additional pointers to blocks of pointers
- Examples: UNIX FFS-based file systems, **ext2**, **ext3**



Multi-Level Indexing



■ Comparison to Indexed Allocation

■ Advantage:

- Does **not waste** space for **unneeded pointers**
- Still fast access for **small files**
- Can grow to what size??

■ Disadvantage:

- Need to read **indirect blocks** of pointers to calculate addresses (extra disk read)
- Keep indirect blocks cached in main memory



Flexible # of Extents

■ Dynamic multiple contiguous regions (**extents**) per file

- Organize **extents** into **multi-level tree structure**
 - Each leaf node: **starting block** and **contiguous size**
 - Minimizes meta-data overhead when have few extents
 - Allows growth beyond fixed number of extents

Fragmentation (internal and external)?

+ Both reasonable

Ability to grow file over time?

+ Can grow

Seek cost for sequential accesses?

+ Still good performance

Speed to calculate random accesses?

+/- Some calculations depending on size

Wasted space for meta-data?

+ Relatively small overhead

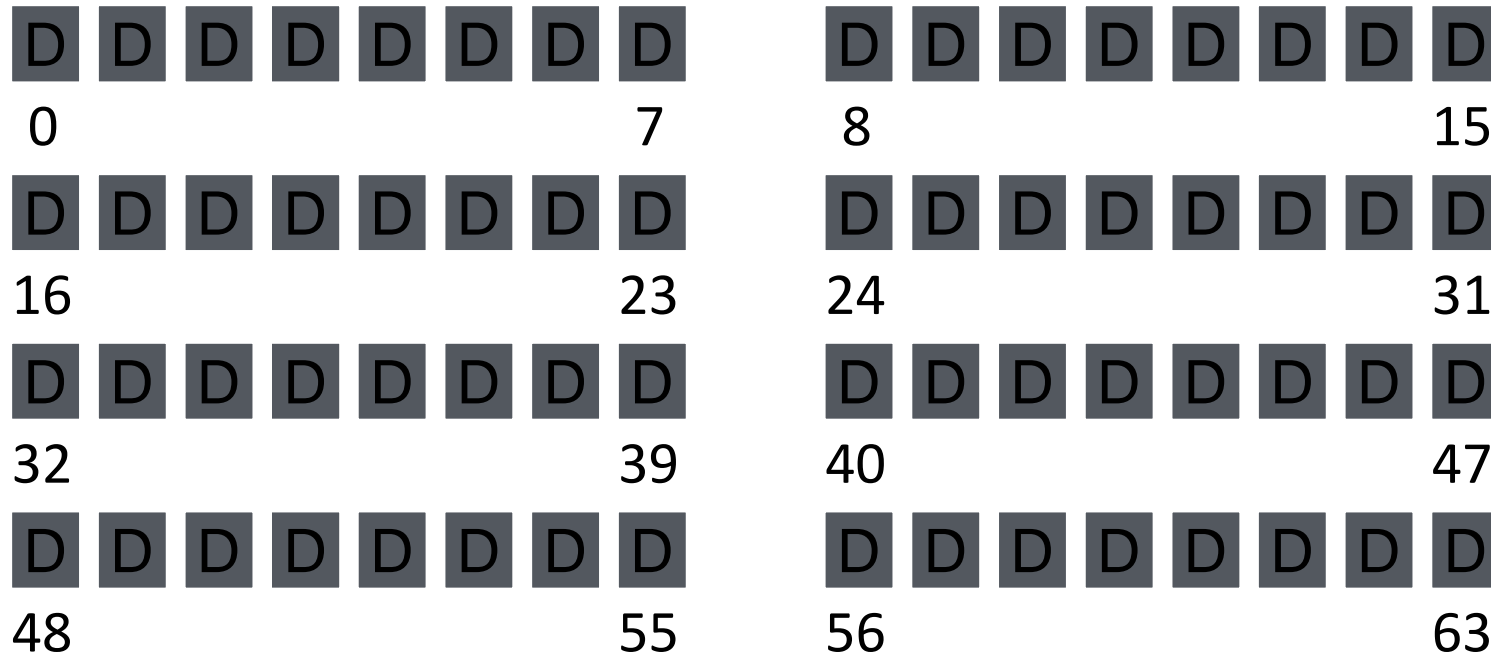
Assume Multi-Level Indexing

- Simple approach
- More complex file systems build from these basic data structures

On-Disk Structures



- data block
- inode table
- indirect block
- directories
- data bitmap
- inode bitmap
- superblock

FS Structs: Empty Disk





Assume each block is 4KB

Data Blocks



Not actual layout : Examine better layout in next lecture
Purpose: Relative number of each time of block

Inodes



One Inode Block


- Each inode is typically 256 bytes (depends on the FS, maybe 128 bytes)
- 4KB disk block
- 16 inodes per inode block.



Inode

type (file or dir?)
uid (owner)
rwX (permissions)
size (in bytes)
Blocks
time (access)
ctime (create)
links_count (# paths)
addrs[N] (N data blocks)

Inodes



Directories

- File systems vary
- Common design: Store **directory entries** in data blocks
 - Large directories just use **multiple data blocks**
 - Use **bit in inode** to distinguish directories from files
- Various formats could be used
 - lists
 - b-trees

Simple Directory List Example

valid	name	inode
1	.	134
1	..	35
1	foo	80
1	bar	23

unlink("foo")

Allocation

- How do we find **free** data blocks or free inodes?
- Free list
- Bitmaps
- Tradeoffs between data structures

Bitmaps?



0

7



16

23



32

39



48

55



8

15



24

31



40

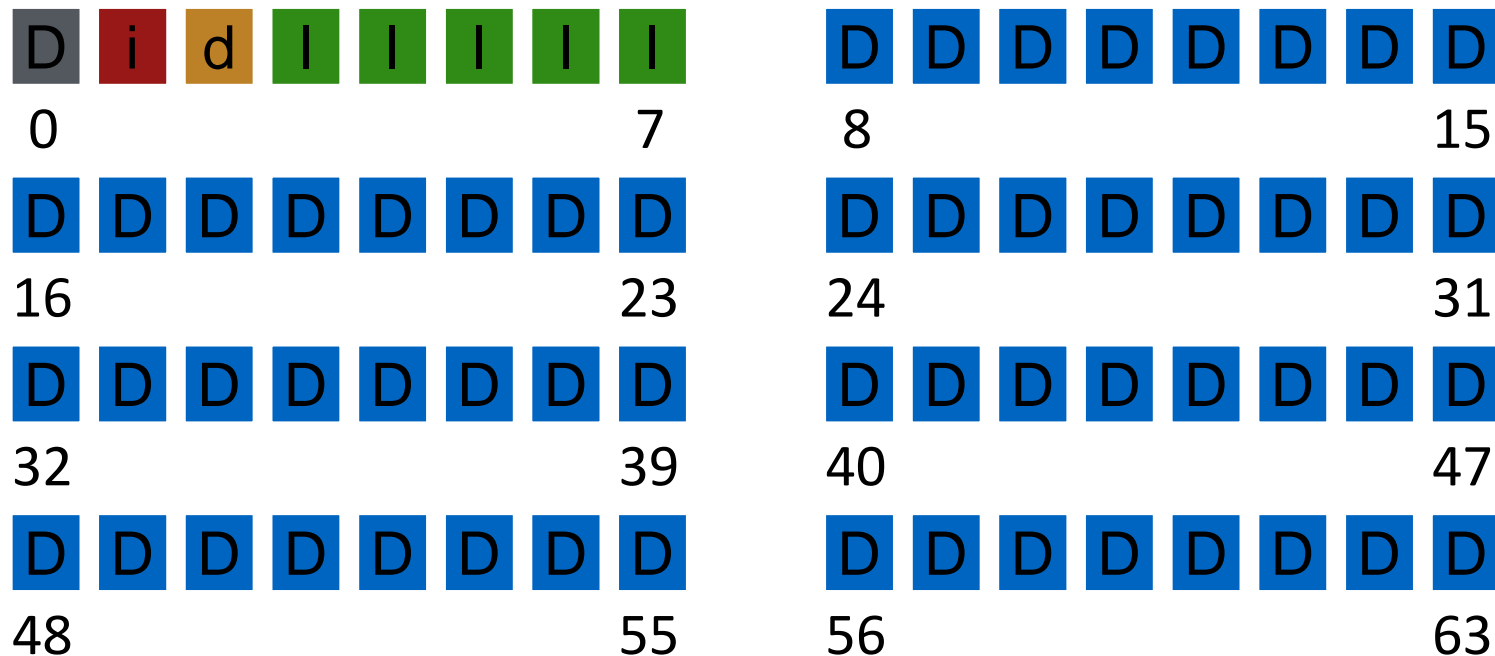
47



56

63

Opportunity for Inconsistency (fsck)



free bitmap for inodes





free bitmap for data blocks

Superblock

- Need to know basic FS configuration metadata, like:
 - block size
 - # of inodes
- Store this in **superblock**

Super Block



super block

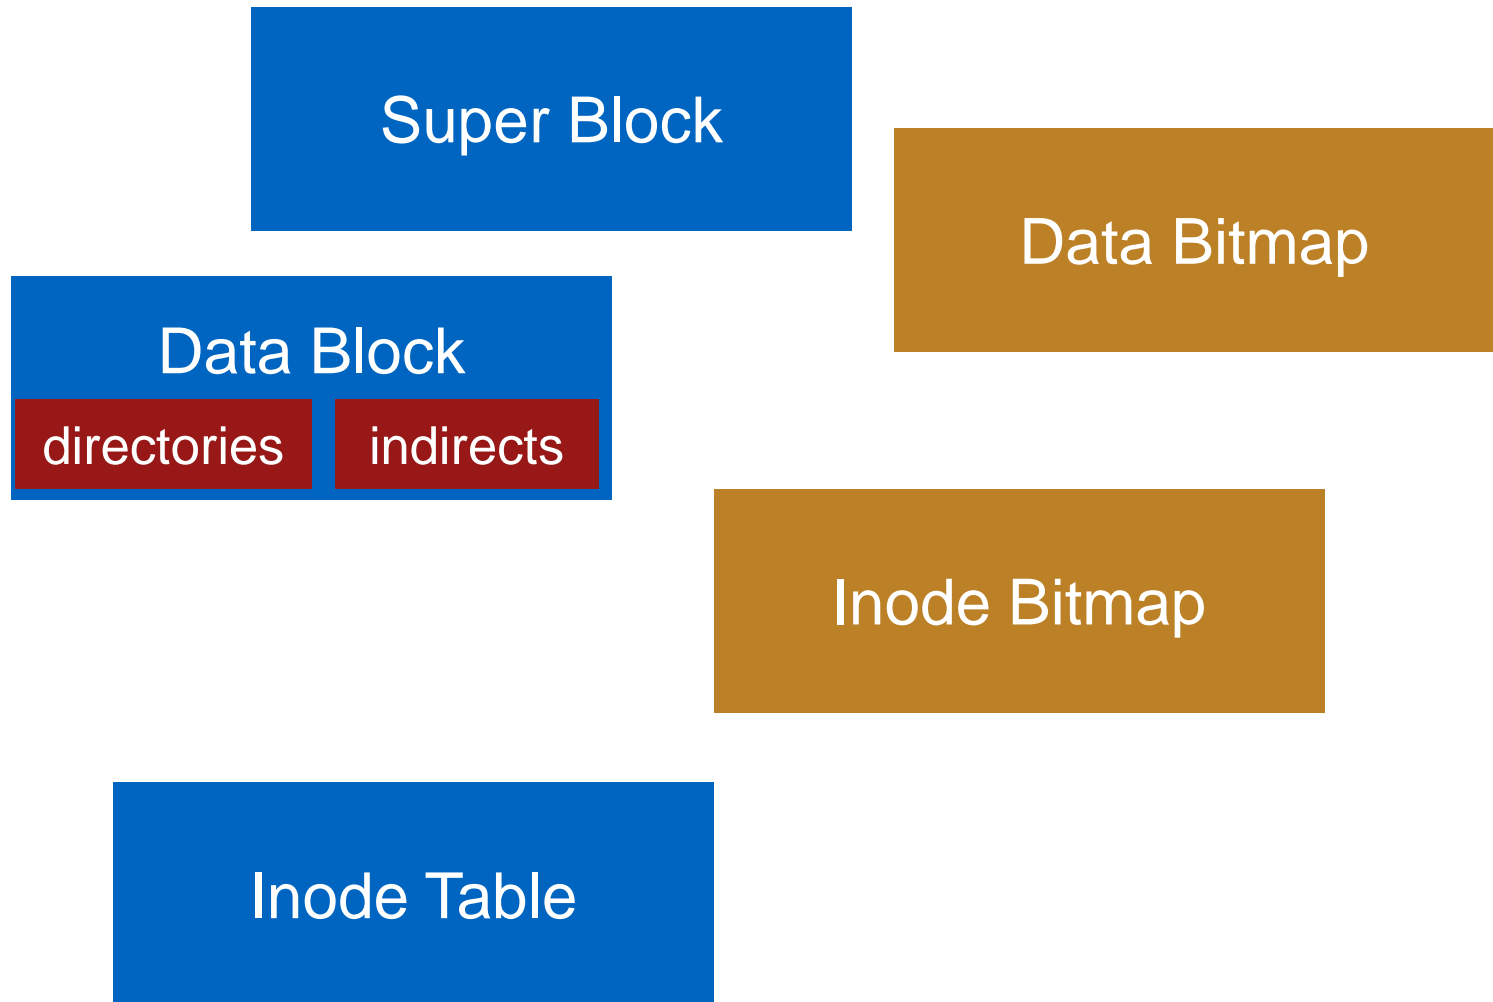


free bitmap for inodes



free bitmap for data blocks

On-Disk Structures



Part 2 : Operations

- create file
- write
- open
- read
- close

create /foo/bar

data bitmap	inode bitmap	root inode	foo/ inode	bar inode	root data	foo/ data
		read			read	
			read			read
	read write					write
				read write		
			write			

What needs to be read and written?

Why read bar inode?

When write partial of a block, it needs to be read from disk
When write an entire block, no read is needed.

open /foo/bar

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
		read			read		
			read			read	
				read			

write to /foo/bar (assume file exists and has been opened)

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
read write				read			write
				write			

read /foo/bar – assume opened

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
				read			
				write			read

Why write bar inode?

Update the access time

close /foo/bar

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data

nothing to do on disk!

Each process minds its business in its file descriptor table

Optimization

- **How can we avoid this excessive I/O for basic ops?**
- **Cache for:**
 - reads
 - write buffering
- **Virtual memory and disk cache**
 - static partitioning – like 10% of memory for disk cache
 - dynamic partitioning – page cache
 - page cache = virtual memory pages + file system pages

Write Buffering

■ Why does procrastination (拖延) help?

- Locality
- Batching
- e.g., an inode with create+update can be batched

■ Overwrites, deletes, scheduling

- Shared structs (e.g., [bitmaps+dirs](#)) often [overwritten](#), reducing the number of writes
- a [temporary](#) file does not need to be write to disk
- batching leads to OS controlled [scheduling](#)

■ We decide: how much to buffer, how long to buffer...

- tradeoffs?
- modern file systems buffer writes between [5-30 seconds](#)

Summary/Future

- **We've described a very simple FS.**
 - basic on-disk structures
 - the basic ops
- **Future questions:**
 - how to allocate efficiently to obtain good performance from disk?
 - how to handle crashes?