

# 统计机器学习 Homework 03

姓名: 雍崔扬

学号: 21307140051

## Problem 1

### Part (1)

Write the Newton-Raphson algorithm to estimate logistic regression, i.e., derive the equation:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta \partial \beta^\top} = - \sum_i x_i x_i^\top p(x_i; \beta) \{1 - p(x_i; \beta)\}$$

**Solution:**

给定训练集  $D_{\text{train}} = \{(x^{(i)}, y_i)\}_{i=1}^n$ ,

用 Logistic 回归模型对每个样本  $x^{(i)}$  进行预测,

输出其标签为 1 的后验概率, 记为  $\hat{y}_i = \sigma(\beta^\top x^{(i)})$  ( $i = 1, \dots, n$ ).

而真实条件概率可以表示为  $\begin{cases} p_r(y_i = 1|x^{(i)}) = y_i \\ p_r(y_i = 0|x^{(i)}) = 1 - y_i \end{cases}$

使用交叉熵损失函数: (简单起见, 不考虑正则化项)

$$\begin{aligned} l(\beta) &= - \sum_{i=1}^n \{p_r(y_i = 1|x) \log(\hat{y}_i) + p_r(y_i = 0|x) \log(1 - \hat{y}_i)\} \\ &= - \sum_{i=1}^n \{y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)\} \\ &\quad (\text{where } \hat{y}_i = \sigma(\beta^\top x^{(i)}) \text{ } (i = 1, \dots, n)) \end{aligned}$$

则我们有:

$$\begin{aligned}
\frac{\partial}{\partial \beta^\top} l(\beta) &= - \sum_{i=1}^n \{y_i \frac{\partial}{\partial \beta^\top} \log(\hat{y}_i) + (1 - y_i) \frac{\partial}{\partial \beta^\top} \log(1 - \hat{y}_i)\} \\
&= - \sum_{i=1}^n \{y_i \frac{1}{\hat{y}_i} \frac{\partial}{\partial \beta^\top} \hat{y}_i + (1 - y_i) \left(-\frac{1}{1 - \hat{y}_i}\right) \frac{\partial}{\partial \beta^\top} \hat{y}_i\} \quad (\text{note that } \frac{\partial}{\partial \beta} \hat{y}_i = \hat{y}_i(1 - \hat{y}_i)x^{(i)}) \\
&= - \sum_{i=1}^n \{y_i \frac{\hat{y}_i(1 - \hat{y}_i)}{\hat{y}_i} x^{(i)} - (1 - y_i) \frac{\hat{y}_i(1 - \hat{y}_i)}{1 - \hat{y}_i} x^{(i)}\} \\
&= - \sum_{i=1}^n \{y_i(1 - \hat{y}_i)x^{(i)} - (1 - y_i)\hat{y}_i x^{(i)}\} \\
&= - \sum_{i=1}^n (y_i - \hat{y}_i)x^{(i)} \\
&\quad (\text{where } \hat{y}_i = \sigma(\beta^\top x^{(i)}) \quad (i = 1, \dots, n)) \\
&= -X^\top(y - \hat{y}) \\
&\quad (\text{where } X = [x^{(1)}, \dots, x^{(n)}]^\top, \hat{y} = [\hat{y}_1, \dots, \hat{y}_n]^\top = \sigma(X\beta))
\end{aligned}$$

$$\begin{aligned}
\frac{\partial^2}{\partial \beta \partial \beta^\top} l(\beta) &= - \sum_{i=1}^n x^{(i)} \frac{\partial}{\partial \beta} (y_i - \hat{y}_i) \quad (\text{note that } \frac{\partial}{\partial \beta} \hat{y}_i = \hat{y}_i^{(i)}(1 - \hat{y}_i^{(i)})(x^{(i)})^\top) \\
&= - \sum_{i=1}^n x^{(i)} \{-\hat{y}_i(1 - \hat{y}_i)(x^{(i)})^\top\} \\
&= \sum_{i=1}^n \hat{y}_i(1 - \hat{y}_i)x^{(i)}(x^{(i)})^\top \\
&\quad (\text{where } \hat{y}_i = \sigma(\beta^\top x^{(i)}) \quad (i = 1, \dots, n)) \\
&= \frac{1}{n} X^\top \text{diag}\{\hat{y} \odot (1_n - \hat{y})\} X \\
&\quad (\text{where } X = [x^{(1)}, \dots, x^{(n)}]^\top, \hat{y} = [\hat{y}_1, \dots, \hat{y}_n]^\top = \sigma(X\beta))
\end{aligned}$$

这样我们就证明了：

$$\frac{\partial^2}{\partial \beta \partial \beta^\top} l(\beta) = \sum_{i=1}^n \hat{y}_i(1 - \hat{y}_i)x^{(i)}(x^{(i)})^\top = \sum_{i=1}^n p(x^{(i)}; \beta)(1 - p(x^{(i)}; \beta))x^{(i)}(x^{(i)})^\top$$

Generate  $X = (1, X_1, X_2)$ , where  $X_j \sim N(0, I_N)$  ( $j = 1, 2$ ).

Set true parameter  $\beta = (0.5, 1.2, -1)^\top$ .

Set  $N = 200, 500, 800, 1000$ .

Estimate  $\beta$  using the Newton-Raphson (NR) algorithm for  $R = 200$  rounds of simulation.

For each round of simulation, terminate the iteration when  $\max_j |\beta_j^{\text{old}} - \beta_j^{\text{new}}| < 10^{-5}$

Denote  $\hat{\beta}_j^{(r)}$  as the estimation of  $\beta_j$  in the  $r$ -th round of simulation.

Then please, for each  $j$ , draw  $\hat{\beta}_j^{(r)} - \beta_j$  in boxplots for  $N = 200, 500, 800, 1000$ .

**Solution:**

生成样本的 python 代码：

```

# Function to generate data
def generate_data(N, true_beta):
    x1 = np.random.normal(size=N)
    x2 = np.random.normal(size=N)
    x = np.column_stack((np.ones(N), x1, x2)) # Adding intercept
    p = 1 / (1 + np.exp(-x @ true_beta)) # Compute probabilities
    y = np.random.binomial(1, p) # Generate binary outcomes
    return x, y

```

纯 Newton 法的 python 代码:

```
# Newton function returning the history of beta
def newton(X, y, max_iter=20, tolerance=1e-5):
    beta = np.zeros(X.shape[1]) # Initialize beta as a 1D array
    for i in range(max_iter):
        p_hat = 1 / (1 + np.exp(-X @ beta)) # Compute predicted probabilities
        gradient = X.T @ (y - p_hat) # Compute the gradient
        hessian = -X.T @ np.diag(p_hat * (1 - p_hat)) @ X # Compute the Hessian
        beta_new = beta - np.linalg.solve(hessian, gradient) # Update beta

        if np.max(np.abs(beta_new - beta)) < tolerance: # Check for convergence
            beta = beta_new
            break

    beta = beta_new # Update beta for next iteration
    return beta
```

函数调用及绘制箱线图的代码:

```
# Set parameters
np.random.seed(51)
true_beta = np.array([0.5, 1.2, -1])
N_values = [200, 500, 800, 1000]
Rounds = 200

# Run simulations
estimates = []
for N in N_values:
    rounds_estimates = []
    for r in range(Rounds):
        X, y = generate_data(N, true_beta)
        rounds_estimates.append(newton(X, y)) # Get the final estimate

    # Print progress every 10%
    if (r + 1) % (Rounds // 10) == 0:
        print(f"Progress for N={N}: Completed {((r + 1) / Rounds) * 100:.1f}%")

    estimates.append(np.array(rounds_estimates))

# Prepare data for boxplot
results = []
for i, N in enumerate(N_values):
    for r in range(Rounds):
        results.append({
            'N': N,
            'beta_hat1': estimates[i][r][0], # Directly use estimates
            'beta_hat2': estimates[i][r][1],
            'beta_hat3': estimates[i][r][2]
        })

results_df = pd.DataFrame(results)

# Melt results for seaborn
results_melted = results_df.melt(id_vars='N', value_vars=['beta_hat1',
    'beta_hat2', 'beta_hat3'],
```

```

var_name='variable', value_name='value')

# Plot boxplots
plt.figure(figsize=(10, 6))
sns.boxplot(x='N', y='value', hue='variable', data=results_melted,
            palette={"beta_hat1": "blue", "beta_hat2": "orange", "beta_hat3":
"green"})
plt.xlabel("Sample Size (N)")
plt.ylabel("Estimation Error ( $\hat{\beta} - \beta$ )")
plt.title("Boxplots of Estimation Errors for Different Sample Sizes")
handles = [
    mpatches.Patch(color='blue', label='β1'),
    mpatches.Patch(color='orange', label='β2'),
    mpatches.Patch(color='green', label='β3')
]
plt.legend(handles=handles, title='')

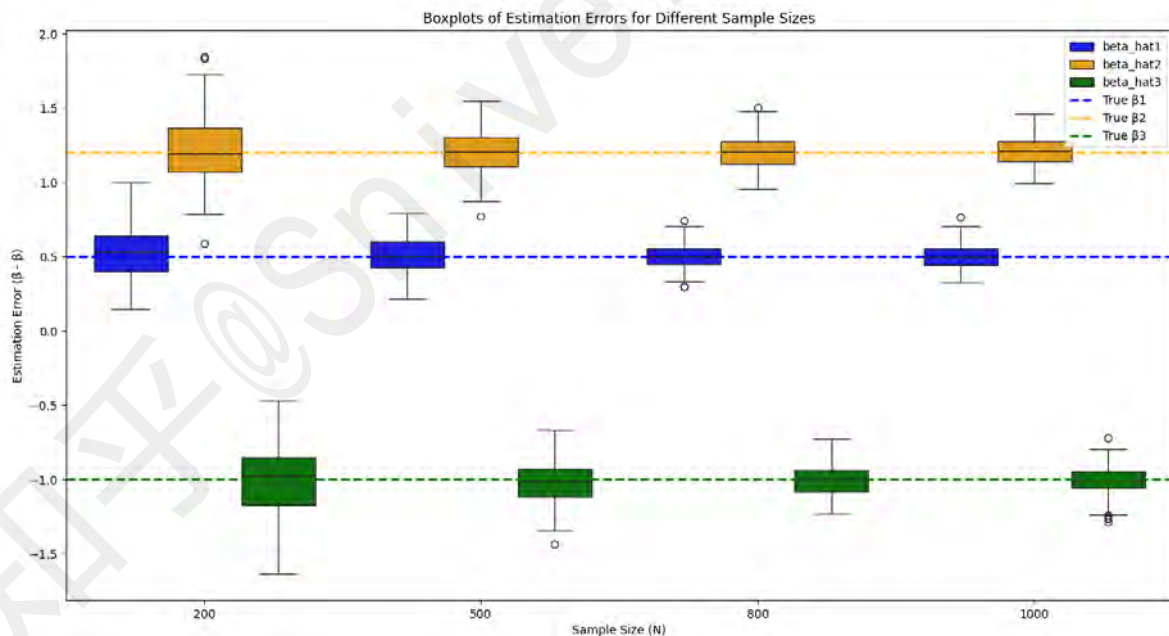
# Plot true values of β
for i, beta in enumerate(true_beta):
    plt.axhline(y=beta, color=['blue', 'orange', 'green'][i], linestyle='--',
linewidth=2, label=f'True β{i+1}')

# Show legend for true values
plt.legend(title='', loc='upper right')

plt.tight_layout()
plt.show()

```

运行结果:



## Part (2)

假设有  $N_+$  个正例和  $N_-$  个负例, 令  $D_+$  与  $D_-$  分别表示正例、负例集合。  
定义排序 "损失" 如下:

$$\ell_{\text{rank}} = \frac{1}{N_+ N_-} \sum_{x^+ \in D_+} \sum_{x^- \in D_-} \left( I(f(x^+) < f(x^-)) + \frac{1}{2} I(f(x^+) = f(x^-)) \right)$$

理解: 若正例的预测值小于负例, 则记 1 个 "罚分", 若相等, 则记 0.5 个罚分.

定义  $AUC := 1 - \ell_{\text{rank}}$

考虑一种简单的情况, 即当数据中不存在  $f(x^+) = f(x^-)$  时, 定义排序 "损失" 如下:

$$\ell_{\text{rank}} = \frac{1}{N_+ N_-} \sum_{x^+ \in D_+} \sum_{x^- \in D_-} I(f(x^+) < f(x^-))$$

试证明以上定义的 AUC 即有限样本的 ROC 曲线下方的面积.

**Proof:**

**(ROC 曲线和 AUC)**

ROC 曲线的纵轴为  $TPR = \text{Recall} = \frac{TP}{TP+FN}$ , 横轴为  $FPR := \frac{FP}{TN+FP}$

而 AUC 是 **ROC 曲线下的面积** (Area Under Curve), 其取值范围为  $[0, 1]$

有限样本集上的 ROC 曲线:

- ① 设分类阈值为 1, 则  $(FPR, TPR) = (0, 0)$
- ② 将预测值从高到低排序, 将阈值依次设为预测值, 分别计算  $(FPR, TPR)$  (可以利用已经计算的结果快速更新)
  - 若当前样本是真正例, 则向纵轴方向走 (FPR 部分不变, TPR 部分增长一个单位  $\frac{1}{N_+}$ )
  - 若当前样本是假正例, 则向横轴方向走 (FPR 部分增长一个单位  $\frac{1}{N_-}$ , TPR 部分不变)

记  $N := N_+ + N_-$

不失一般性, 假设预测值是从高到低排序的 (即有  $f(x_1) \geq \dots \geq f(x_N)$ )

其中负例的指标为  $i_1, \dots, i_{N_-}$

则我们有:

$$\begin{aligned} AUC &= 1 - \ell_{\text{rank}} \\ &= 1 - \frac{1}{N_- N_+} \sum_{x_- \in D_-} \sum_{x^+ \in D_+} [1 - I(f(x_+) < f(x_-))] \\ &= \frac{1}{N_- N_+} \sum_{x_- \in D_-} \sum_{x^+ \in D_+} I(f(x_+) \geq f(x_-)) \\ &= \frac{1}{N_- N_+} \sum_{x_- \in D_-} \text{Number of positive samples } x_+ \text{ that satisfies } f(x_+) \geq f(x_-) \\ &= \frac{1}{N_- N_+} \sum_{j=1}^{N_-} \text{Number of positive samples in } \{x_1, \dots, x_{i_j}\} \\ &= \sum_{j=1}^{N_-} \frac{1}{N_-} \frac{i_j - j}{N_+} \\ &= S(\text{Area Under ROC Curve}) \end{aligned}$$

因此  $AUC := 1 - \ell_{\text{rank}}$  即为 ROC 曲线下方的面积.

## Problem 2

客户流失预警数据:

- 训练数据集: `sampladata.csv`
- 测试数据集: `preddata.csv`

数据文件来自国内某运营商, 数据已经进行了清理, 数据集共 8 个变量:

	变量名		详细说明	备注
因变量 (下月)	churn	是否流失	1=流失 0=不流失	流失率 1.25%
自变量 (当月)	tenure	在网时长	连续变量 单位: 天	客户从入网到截止数据提取日期时在网时间
	expense	当月花费	连续变量 单位: 元	客户在提取月份时的花费总额
	degree	个体的度	连续变量 单位: 人数	和客户通话的总人数, 去重之后的呼入与呼出加总
	tightness	联系强度	连续变量 分钟/人	通话总时间除以总人数
	entropy	个体信息熵	连续变量	$E_i = -\sum_{a_{ij}=1} p_{ij} * \log(p_{ij})$ , 其中 $E_i$ 为个体 i 的信息熵, $a_{ij} = 1$ 代表个体 i 和 j 通过电话, $p_{ij}$ 代表 j 和 i 通话的分钟数据占 i 总通话分钟的比例
	chgdegree	个体度的变化	连续变量 单位: %	(本月个体的度-上月个体的度)/上月个体的度
	chgexpense	花费的变化	连续变量 单位: %	(本月花费-上月花费)/上月花费

## Part (1)

读入数据:

```
library(ROSE)

# 读取训练数据集并过采样
train_data <- read.csv("sampledata.csv")
train_data <- ovun.sample(churn ~ ., data = train_data, method = "over", N = 2 *
nrow(train_data))$data

# 读取测试数据集并过采样
test_data <- read.csv("preddata.csv")
test_data <- ovun.sample(churn ~ ., data = test_data, method = "over", N = 2 *
nrow(test_data))$data
```

其中过采样是因为我们发现在训练集中, 正例只占 1.24%

```
> train_data <- read.csv("sampledata.csv")
> sum(train_data$churn == 1)
[1] 602
> nrow(train_data)
[1] 48285
> 602 / 48285 * 100
[1] 1.246764
```

## Part (2)

绘制因变量和各个自变量的箱线图 (提示: 可以对右偏分布的数据取对数)

**Solution:**

```
library(ggplot2)
library(reshape2)
```

```

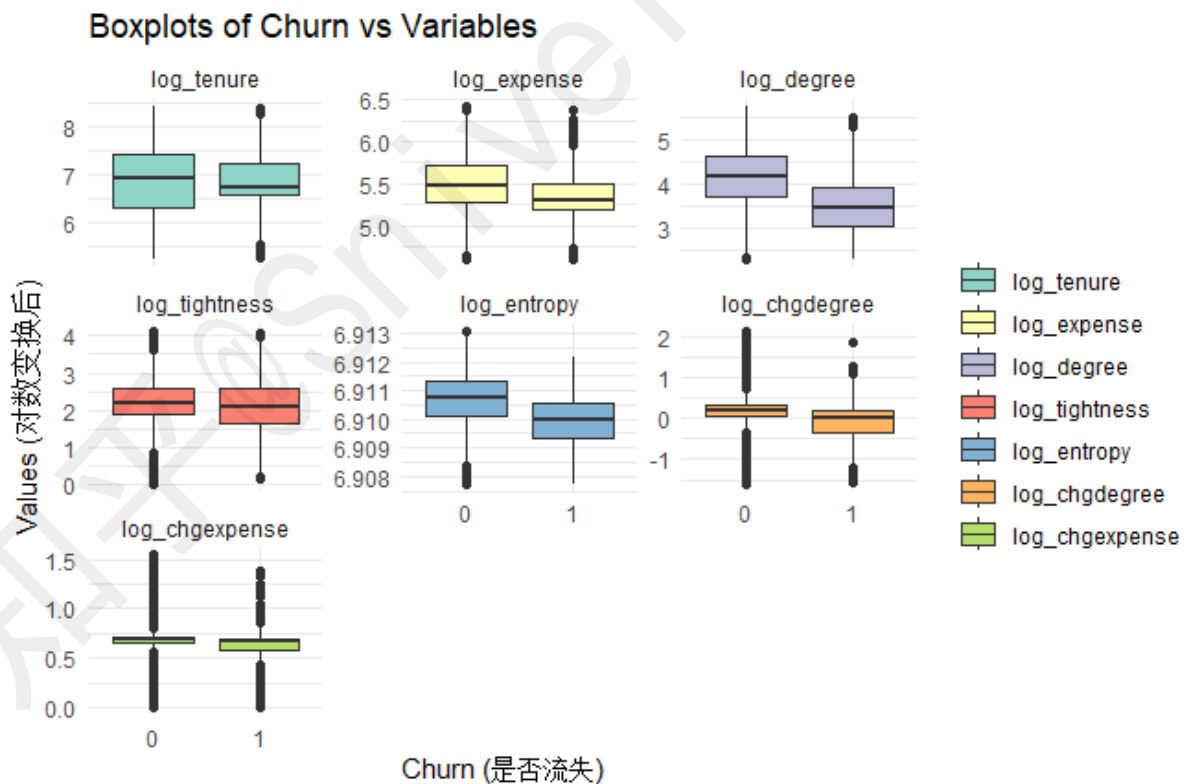
# 对右偏分布的数据进行对数变换（以防止负值）
alpha = c(0, 1e2, 1e1, 1, 1e3, 1.2, 2)
train_data$log_tenure <- log(train_data$tenure + alpha[1])
train_data$log_expense <- log(train_data$expense + alpha[2])
train_data$log_degree <- log(train_data$degree + alpha[3])
train_data$log_tightness <- log(train_data$tightness + alpha[4])
train_data$log_entropy <- log(train_data$entropy + alpha[5])
train_data$log_chgdegree <- log(train_data$chgdegree + alpha[6])
train_data$log_chgexpense <- log(train_data$chgexpense + alpha[7])

# 创建一个包含所有自变量和因变量的数据框
melted_data <- melt(train_data, id.vars = "churn",
                    measure.vars = c("log_tenure", "log_expense", "log_degree",
                                     "log_tightness", "log_entropy",
                                     "log_chgdegree", "log_chgexpense"))

# 绘制箱线图
ggplot(melted_data, aes(x = factor(churn), y = value, fill = variable)) +
  geom_boxplot() +
  labs(x = "Churn (是否流失)", y = "Values (对数变换后)",
       title = "Boxplots of Churn vs Variables") +
  theme_minimal() +
  scale_fill_brewer(palette = "Set3") +
  theme(legend.title = element_blank()) +
  facet_wrap(~ variable, scales = "free_y")

```

运行结果:



## Part (3)

以是否流失为因变量，使用 `scale()` 函数对自变量进行标准化 (使其均值为 0，方差为 1)  
使用 `glm()` 函数建立逻辑回归模型，给出系数估计结果，并对结果进行解读。

**Solution:**

```
# 标准化自变量
train_data$standardized_tenure <- scale(train_data$tenure)
train_data$standardized_expense <- scale(train_data$expense)
train_data$standardized_degree <- scale(train_data$degree)
train_data$standardized_tightness <- scale(train_data$tightness)
train_data$standardized_entropy <- scale(train_data$entropy)
train_data$standardized_chgdegree <- scale(train_data$chgdegree)
train_data$standardized_chgexpense <- scale(train_data$chgexpense)

# 建立逻辑回归模型
model <- glm(churn ~ standardized_tenure + standardized_expense +
             standardized_degree + standardized_tightness +
             standardized_entropy + standardized_chgdegree +
             standardized_chgexpense,
             data = train_data,
             family = binomial)

# 输出系数估计结果
summary(model)
```

输出结果:

```
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    0.007507   0.007542   0.995    0.32
standardized_tenure -0.179136  0.008318 -21.535 <2e-16 ***
standardized_expense -0.242420  0.008585 -28.236 <2e-16 ***
standardized_degree -0.467112  0.016561 -28.205 <2e-16 ***
standardized_tightness -0.224234  0.007843 -28.590 <2e-16 ***
standardized_entropy -0.498783  0.015555 -32.066 <2e-16 ***
standardized_chgdegree -0.259828  0.008355 -31.099 <2e-16 ***
standardized_chgexpense -0.122920  0.007941 -15.478 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 133859 on 96569 degrees of freedom
Residual deviance: 111191 on 96562 degrees of freedom
AIC: 111207

Number of Fisher Scoring iterations: 4
```

结果解读:

- 模型残差偏差 (Residual deviance) 为 111191 (因变量数据的固有偏差为 133859)，AIC 值为 111207  
表明该模型具有较好的拟合效果且复杂度不高。



- \*\*\* 表示  $p$  值小于 0.001, \*\* 表示  $p$  值小于 0.01, \* 表示  $p$  值小于 0.05  
上述结果显示截距项和 7 个自变量均极为显著。

## Part (4)

使用建立好的逻辑回归模型，分别使用 `predict()` 函数对训练集和测试集进行预测，得到每个用户的预测流失概率值。

**Solution:**

```
# 训练集预测
train_data$predicted_probabilities <- predict(model, newdata = train_data, type =
"response")
head(train_data[, c("churn", "predicted_probabilities")]) # 查看训练集的预测结果

# 测试集预测
test_data$standardized_tenure <- scale(test_data$tenure)
test_data$standardized_expense <- scale(test_data$expense)
test_data$standardized_degree <- scale(test_data$degree)
test_data$standardized_tightness <- scale(test_data$tightness)
test_data$standardized_entropy <- scale(test_data$entropy)
test_data$standardized_chgdegree <- scale(test_data$chgdegree)
test_data$standardized_chgexpense <- scale(test_data$chgexpense)

test_data$predicted_probabilities <- predict(model, newdata = test_data, type =
"response")
head(test_data[, c("predicted_probabilities")]) # 查看测试集的预测结果
```

## Part (5)

基于 Part(4) 中预测的结果，分别使用 R 包 `pROC` 中的 `plot.roc()` 绘制训练集和测试集上预测结果的 ROC 曲线，  
计算相应的 AUC 值，并根据 ROC 曲线和 AUC 值对模型进行评价。

**Solution:**

```
library(pROC)

# 绘制训练集的 ROC 曲线
roc_train <- roc(train_data$churn, train_data$predicted_probabilities)
plot(roc_train, main = "ROC Curve for Training Set", col = "blue")
auc_train <- auc(roc_train)
print(paste("AUC on train data:", round(auc_train, 3)))
text(0.6, 0.4, paste("AUC =", round(auc_train, 3)), col = "blue")


# 绘制测试集的 ROC 曲线
roc_test <- roc(test_data$churn, test_data$predicted_probabilities)
plot(roc_test, main = "ROC Curve for Testing Set", col = "red", add = TRUE)
auc_test <- auc(roc_test)
print(paste("AUC on test data:", round(auc_test, 3)))
text(0.6, 0.3, paste("AUC =", round(auc_test, 3)), col = "red")

# 添加图例
legend("bottomright", legend = c("Training AUC", "Testing AUC"),
      col = c("blue", "red"), lwd = 2)
```

运行结果:

"AUC on train data: 0.772"

"AUC on test data: 0.779"



模型无论是在测试集还是训练集上 AUC 值都远离 0.5，接近 1，说明其对顾客是否流失的预测能力较强。(AUC 值越接近于 1，当前的分类算法越有可能将正样本排在负样本前面，即能够更好的分类)

## Part (6)

(自己补充的题目)

基于 Part (5) 绘制的 ROC 曲线，设定最优阈值用于对预测结果进行二分类。

**Solution:**

```
# 计算 Youden's J
coords <- coords(roc_train, "best", ret = c("threshold", "sensitivity",
"specificity"))

# 输出最佳阈值
best_threshold <- as.numeric(coords["threshold"])
print(paste("Best threshold:", round(best_threshold, 3)))

# 二分类：将预测概率转换为分类标签
test_data$predicted_class <- ifelse(test_data$predicted_probabilities >=
best_threshold, 1, 0)

# 计算分类准确率
train_accuracy <- mean(train_data$predicted_class == train_data$churn)
print(paste("训练集分类准确率:", round(train_accuracy * 100, 2), "%"))
test_accuracy <- mean(test_data$predicted_class == test_data$churn)
print(paste("测试集分类准确率:", round(test_accuracy * 100, 2), "%"))
```

输出结果:

"Best threshold: 0.575"

"训练集分类准确率: 70.89 %"

"测试集分类准确率: 58.93 %"

**The End**

知乎@Snivellus Snape